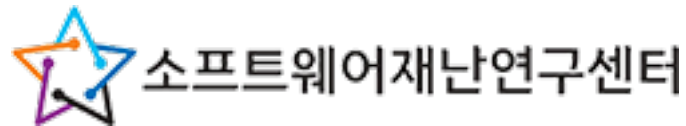


Research on System Call Filtering Technique for Defense against Host Kernel Exploits in Container Environment

Somin Song (Advisor: Byungchul Tak)

Kyungpook National University (KNU), Daegu, Republic of Korea



July 13rd, 2022

Rapid Growth in Cloud Adaption and Security Concern

- The total volume of data that will be stored in the cloud by 2025, which accounts for **50%** of all the data in the world

– ArcServe, 2020 [1]

- According to 74% of global IT decision-makers, **95%** of all workload will be in the cloud within the next five years

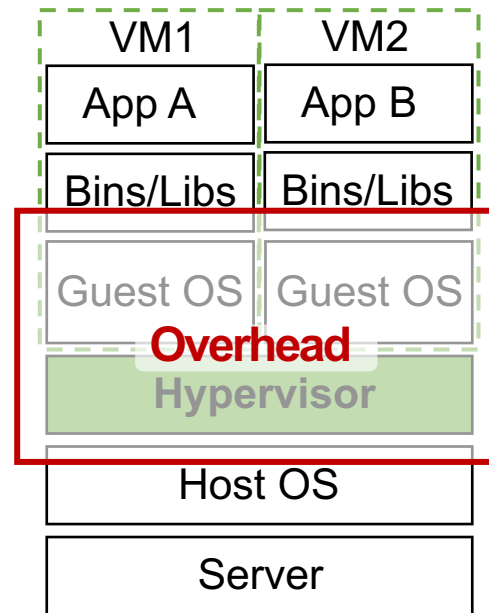
– LogicMonitor, 2020 [1]

- **75%** of enterprises and **90%** of cyber security experts agree that **security is their top concern**

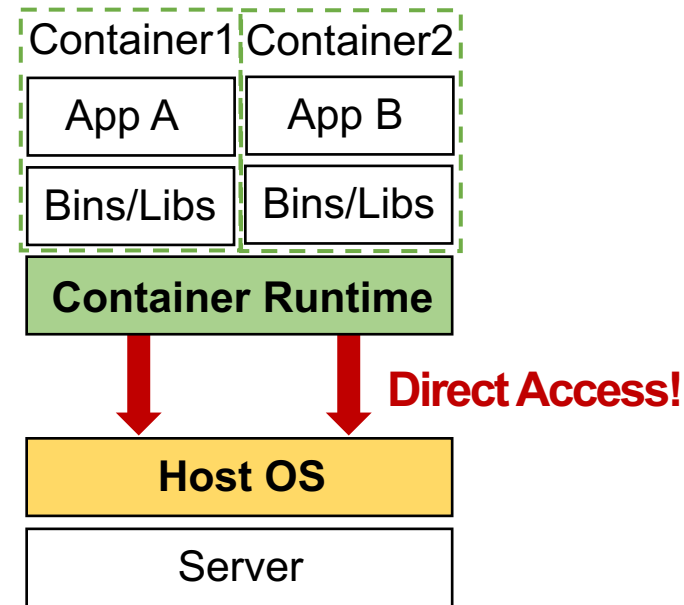
– ArcServe, 2020 [1]

Container Security in Cloud Native Environment

- Containers are a **key virtualization technology** in cloud computing



VM

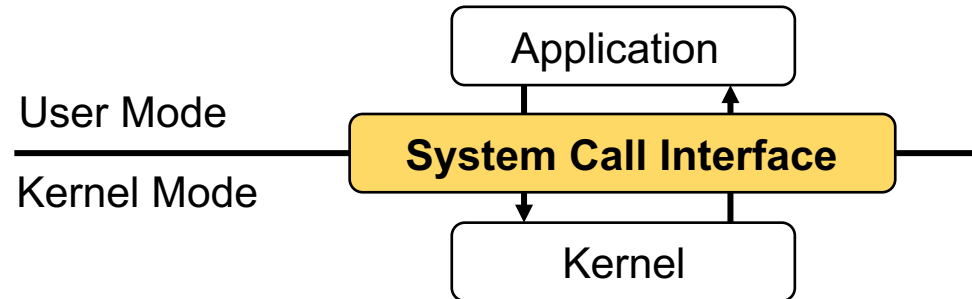


Container

Containers share a Host OS!

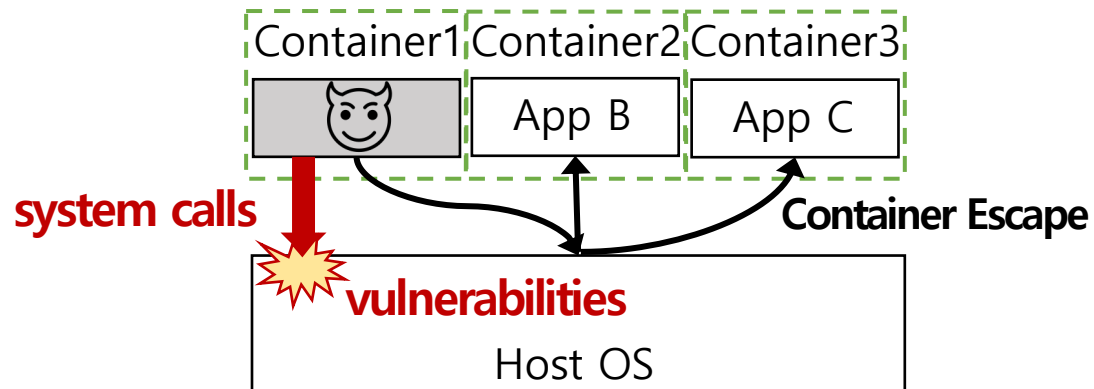
Container Escape through System Call

- System Call



- Container Escape

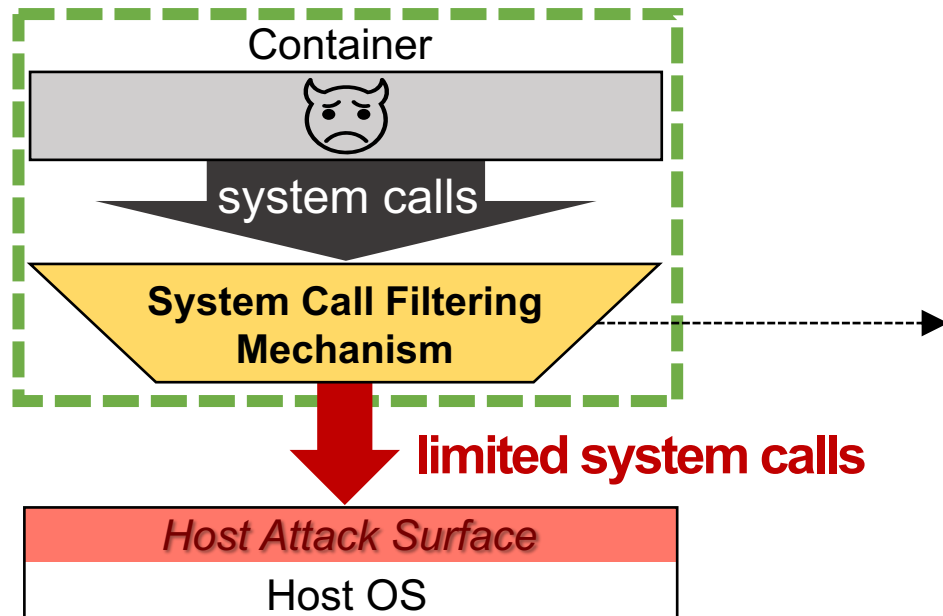
- Exploitation of Host kernel vulnerabilities through carefully **crafted system calls**



- Dirty COW Docker Escape (CVE-2016-5195)
- Runc Container Escape (CVE-2019-5736)
- Kubernetes Container Escape (CVE-2022-0185)
- Dirty Pipe Container Escape (CVE-2022-0847)

System Call Filtering Protection Mechanism

The fewer system calls, the more secure!

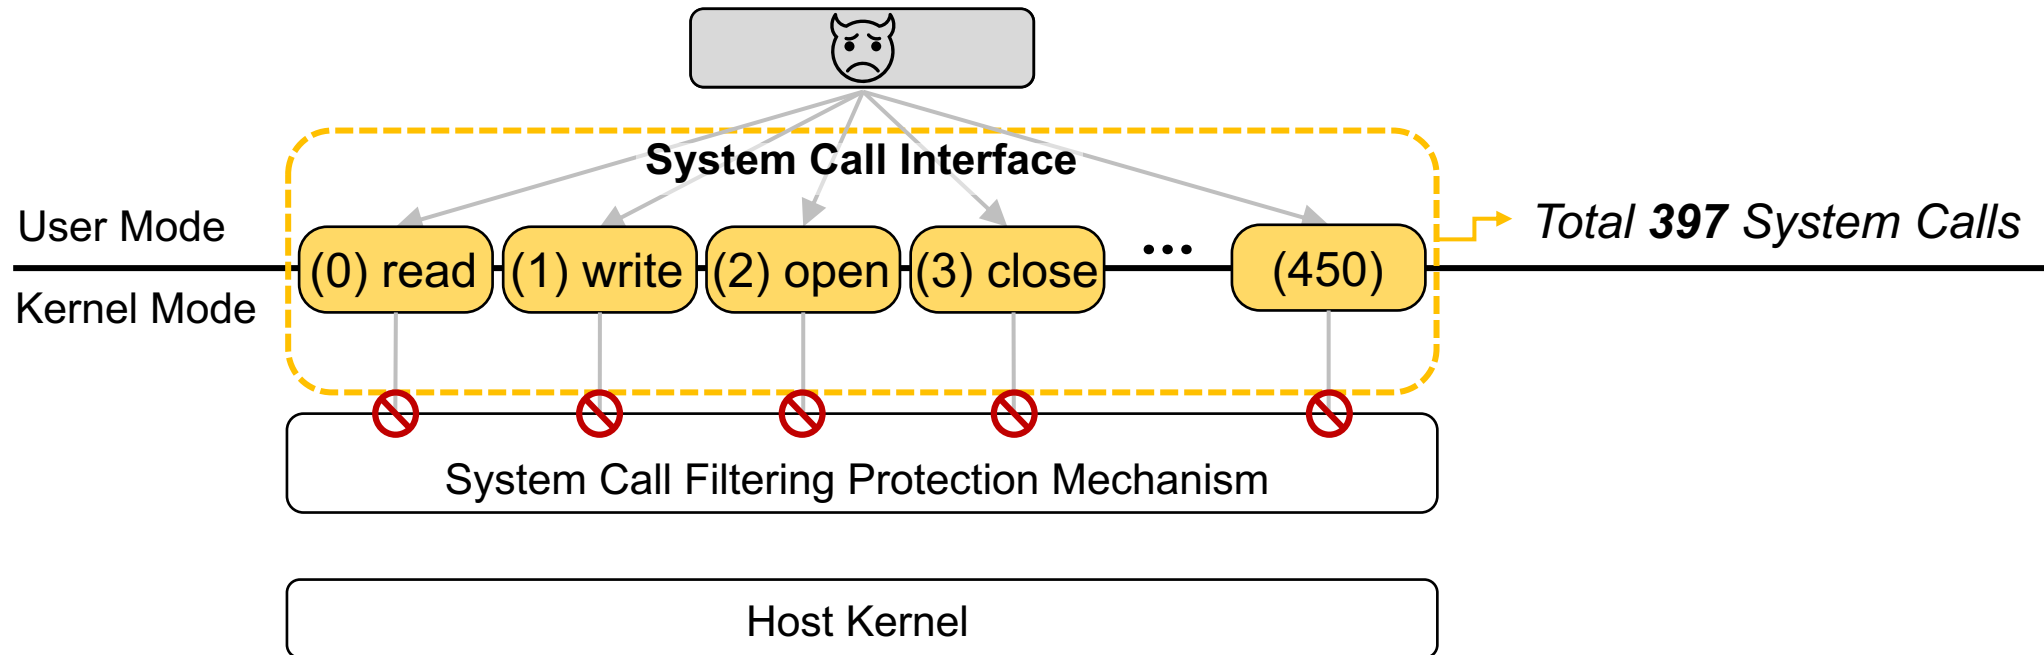


Seccomp (Secure Computing mode)

- Linux kernel feature
- Blocks the system calls
- Restrict a Container's system calls

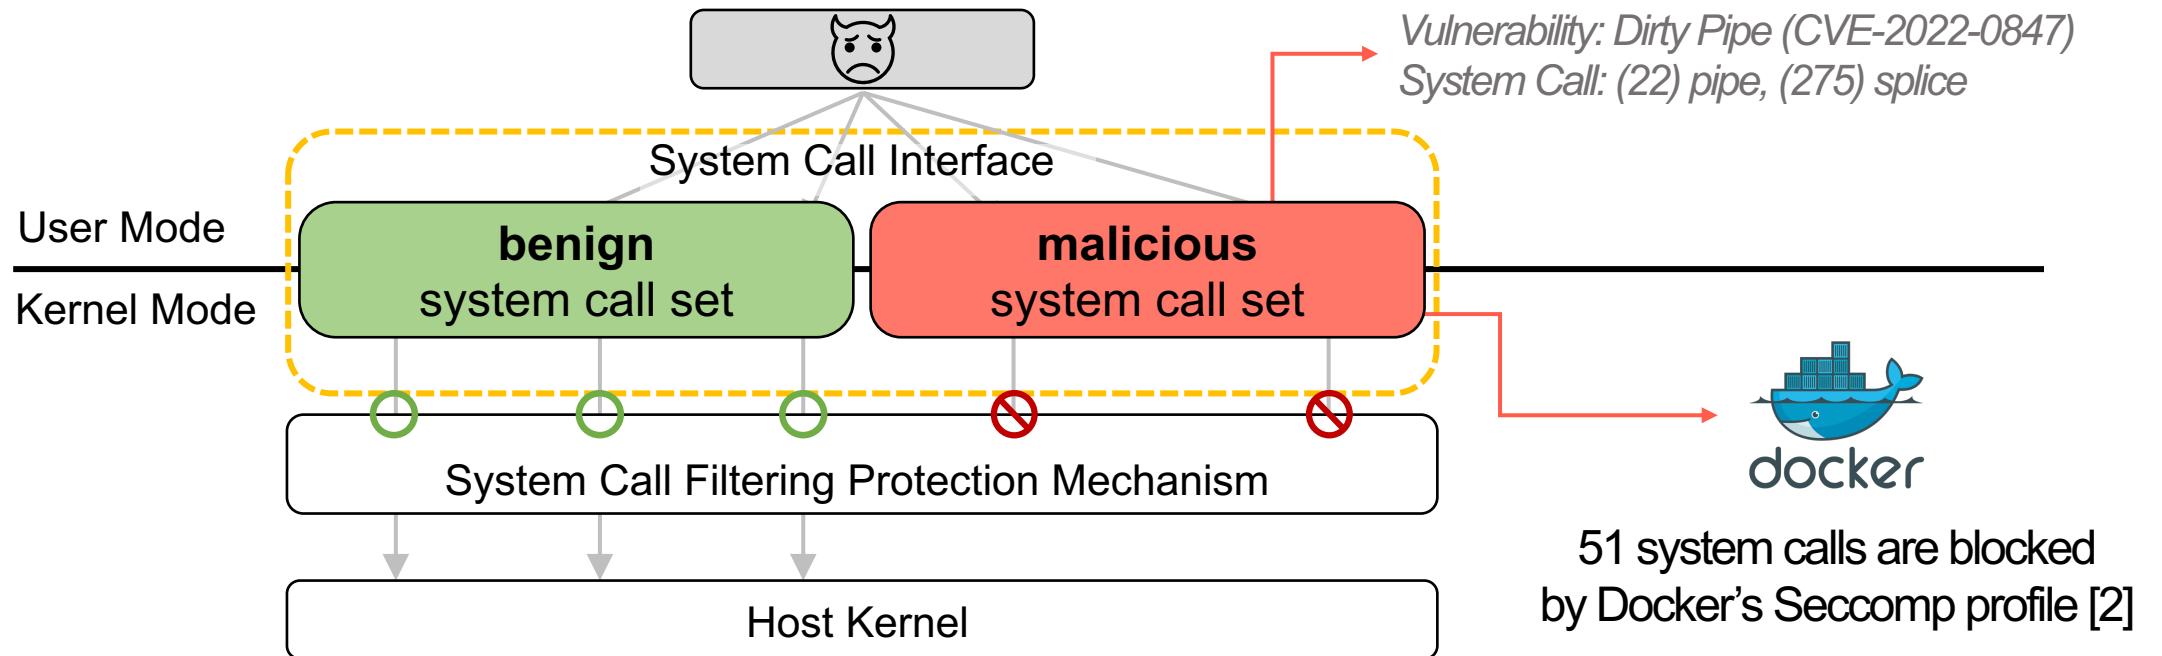
System Call Filtering Protection Mechanism

- “Which System Call should be **allowed/blocked**?”



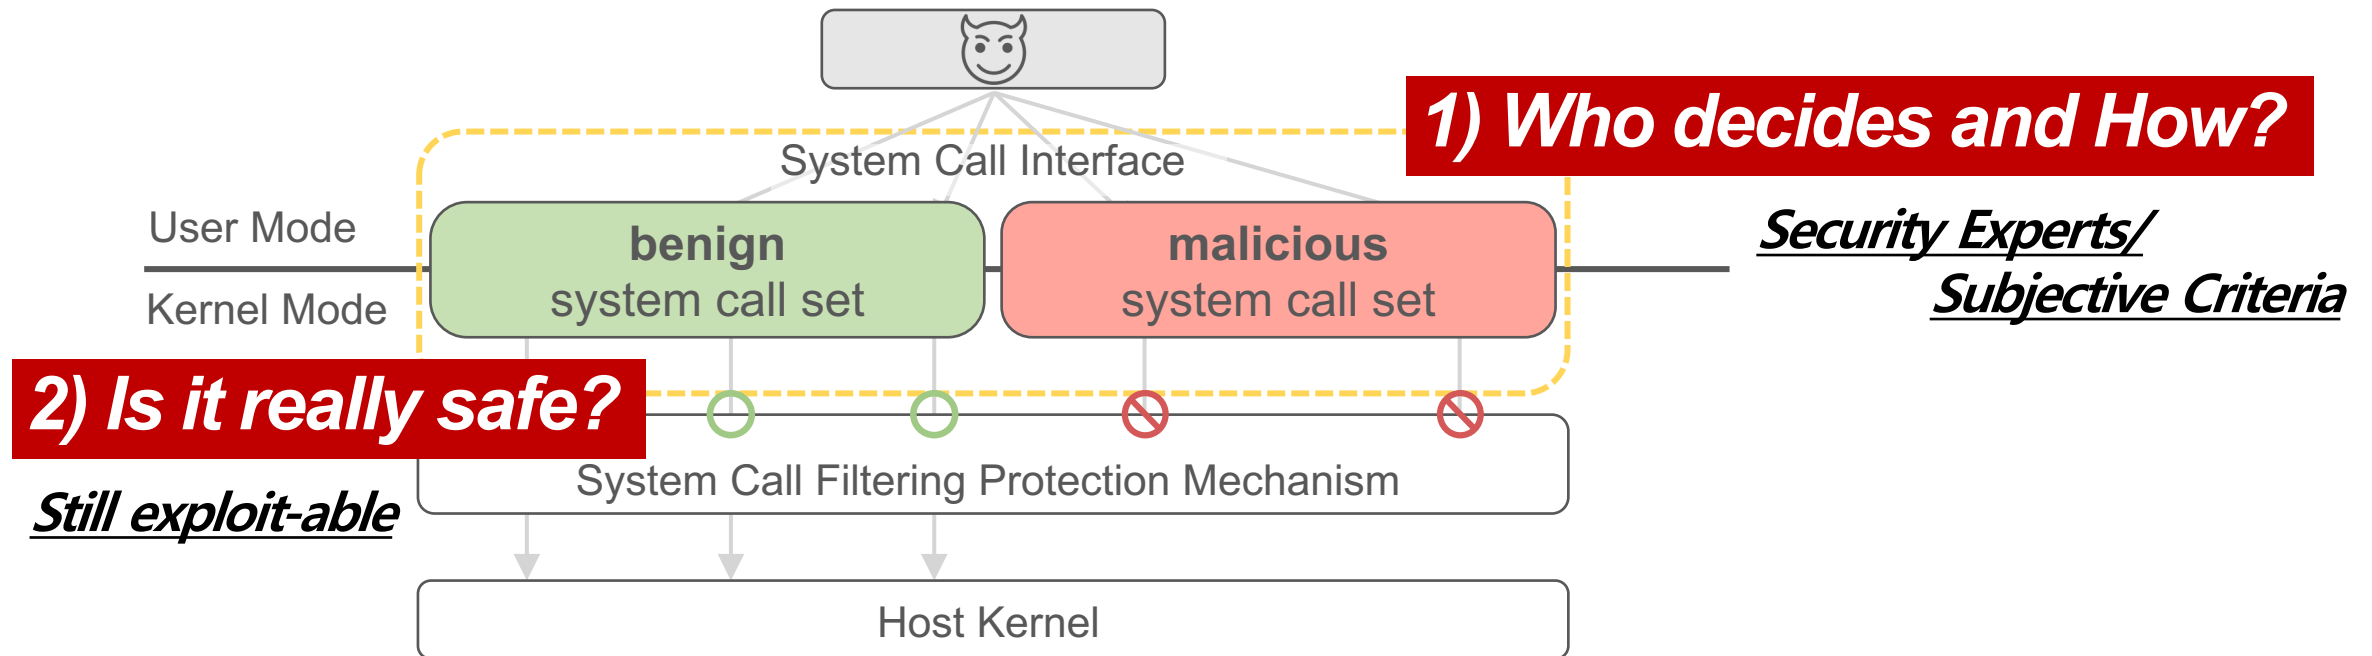
System Call Filtering Protection Mechanism

- “Which System Call should be allowed/blocked?”
 - Which System Call is benign/malicious?



Problems

- “Which System Call should be allowed/blocked?”
 - Which System Call is benign/malicious?



Two Novel Approaches for System Call Filtering

- “Research on System Call Filtering Technique for Defense against Host Kernel Exploits in Container Environment”
 - Goal: To improve the security capability of the system call filtering

Problem 1) Who decides and How?

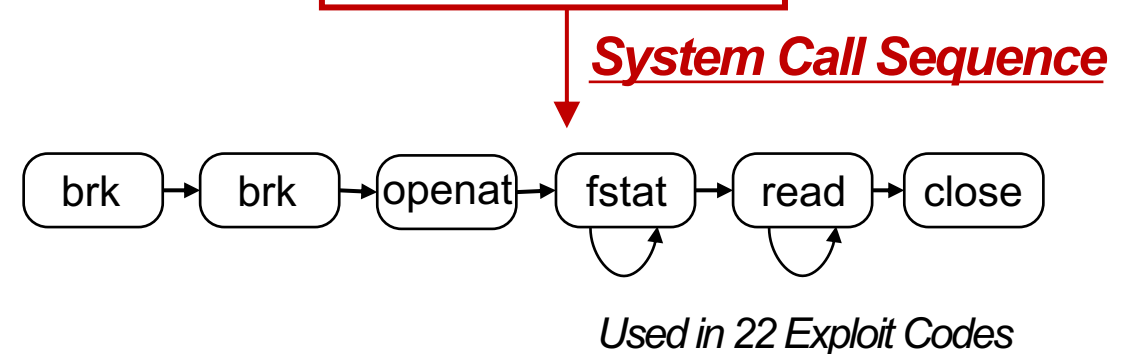
- Limitation of Problem 1
 - Manual analysis by experts through subjective criteria
- Approach 1
 - Risk auto-assessment of system calls using objective criteria

Top Rank capset : 0.49551

Last Rank close : 0.026151

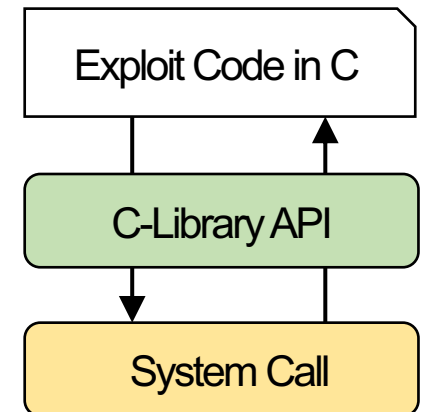
Problem 2) is it really safe?

- Limitation of Problem 2
 - Exploitable by combining allowed system calls
- Approach 2
 - Exploitable system call combination investigation



Key Exploit Code Analysis Methodologies

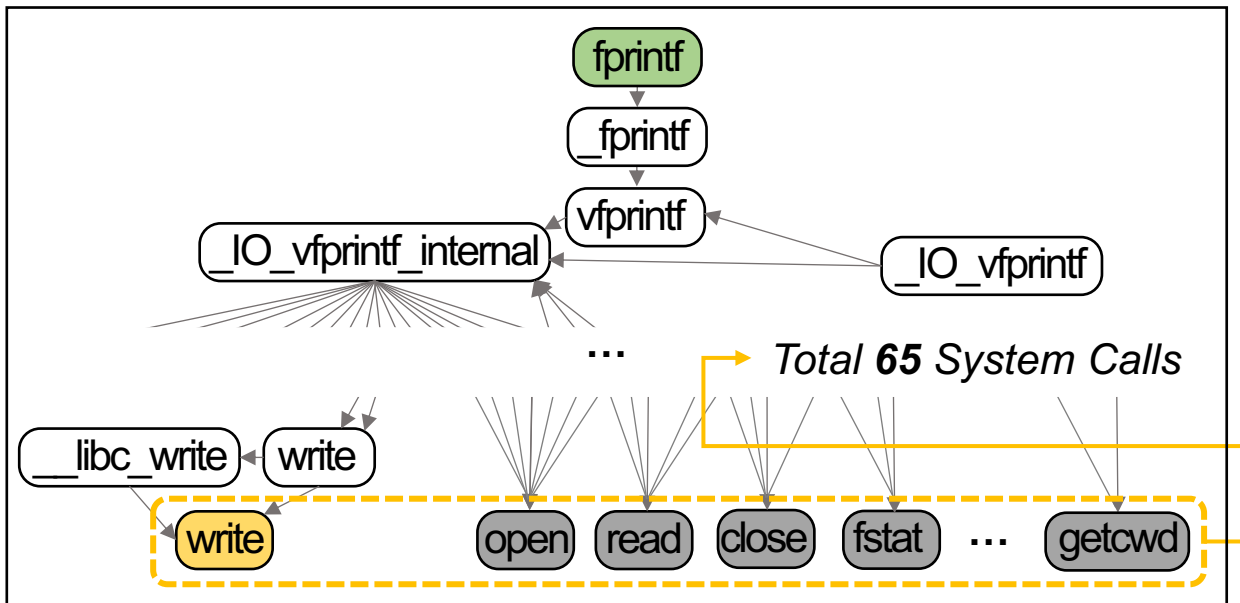
- Goal: To **investigate the system calls** that the exploit codes invoke
 - System Call Set / System Call Sequence
- Only Dynamic Analysis
 - Method of tracing system calls while **directly executing** the exploit codes
 - + Upon successful execution, complete information can be obtained
 - Exploit code is system-dependent
 - ▶ It is very difficult to set up the environment to run the exploit code
 - ▶ The Linux kernel has thousands of version histories
- Only Static Analysis
 - Method of analyzing the source code **without executing** the exploit codes
 - + Automated Large-scale analysis is possible
 - ▶ No configuration is required for the code execution
 - C language makes static analysis difficult
 - ▶ Dynamic linking/allocation and pointers



Problems in Precedent Researches

- Challenge in **Static Analysis**

- Confine [3] (RAID '20) : built libc-to-syscall call graph



- **missing/boating system calls**

- Nicholas et al. [4], Abubakar et al. [5], and Olufogorehan et al. [6] : generated syscall lists
 - **Do not attempt to mine for system call sequences**

- Challenge in **Dynamic Analysis**

- Lopes et al. [7] : used unit testing combined with fuzzing by running the target application
- Wan et al. [8] : created custom profiles with pre-defined test suites
- Speaker [9] : traced system call during a container execution

- **users must have a comprehensive understanding of the applications' behavior**

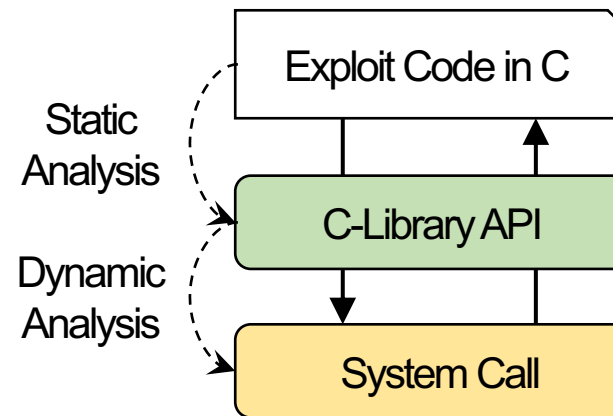
- **Not sound to find all corner cases**

- **Do not attempt to mine for system call sequences**

Key Exploit Code Analysis Methodologies

- Hybrid Analysis

- Combine static analysis and dynamic analysis to generate a system call set/sequence corresponding to each exploit code
- **Static Analysis:** to **extract library function set/sequence** on all possible control flows where the exploits can be successfully triggered
- **Dynamic Analysis:** to build a **mapping between library functions and system call set/sequences**



Static Analysis to Map Exploit Code-Glibc function

- Using GCC + GCC GIMPLE IR (Intermediate Representation)
 - Getting info on [control flow of glibc function](#)

```
void foo(){  
    int i,j;  
    do{  
        i = getuid();  
        j = geteuid();  
    } while(i == j);  
    bar();  
}
```

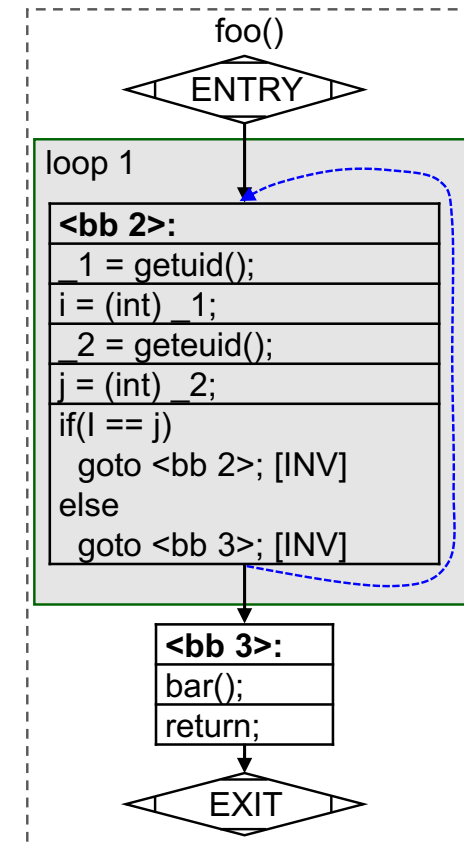
Example source code

```
foo: getuid-geteuid-getuid-geteuid-getuid-geteuid-bar
```

Library function call sequence(s) for code

```
;; Function foo (...)  
;; ...  
;; Loop 1 ...  
;; header 2, latch 2  
;; depth 1, outer 0  
;; nodes: 2  
;; 2 succs { 2 3 }  
;; 3 succs { 1 }  
foo ()  
{  
    ...  
    <bb 2> :  
    _1 = getuid ();  
    i = (int) _1;  
    _2 = geteuid ();  
    j = (int) _2;  
    if (i == j)  
        goto <bb 2>; [INV]  
    else  
        goto <bb 3>; [INV]  
    <bb 3> :  
    bar ();  
    return;  
}
```

GCC's GIMPLE IR



Visualization

Dynamic Analysis to Map Glibc function-System Call

- Dynamic Analysis (using API Sanity Checker dataset [10] + ftrace mechanism)

(a) API Sanity Checker Dataset

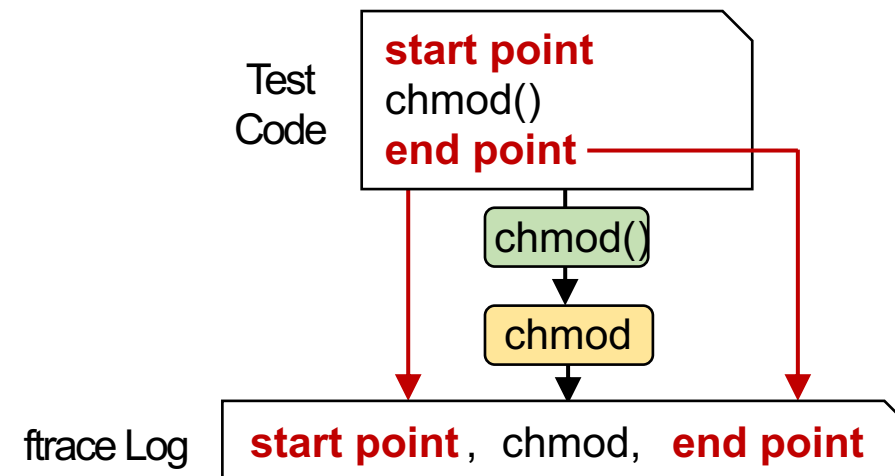
- An automatic generator of basic unit tests for a shared C/C++ library

```
#include <rpc/types.h>
#include <sys/stat.h>

int main(int argc, char *argv[])
{
    __mode_t __mode = umask(0);
    chmod ((const char *) "/proc/self/exec", 3565);
    return 0;
}
```

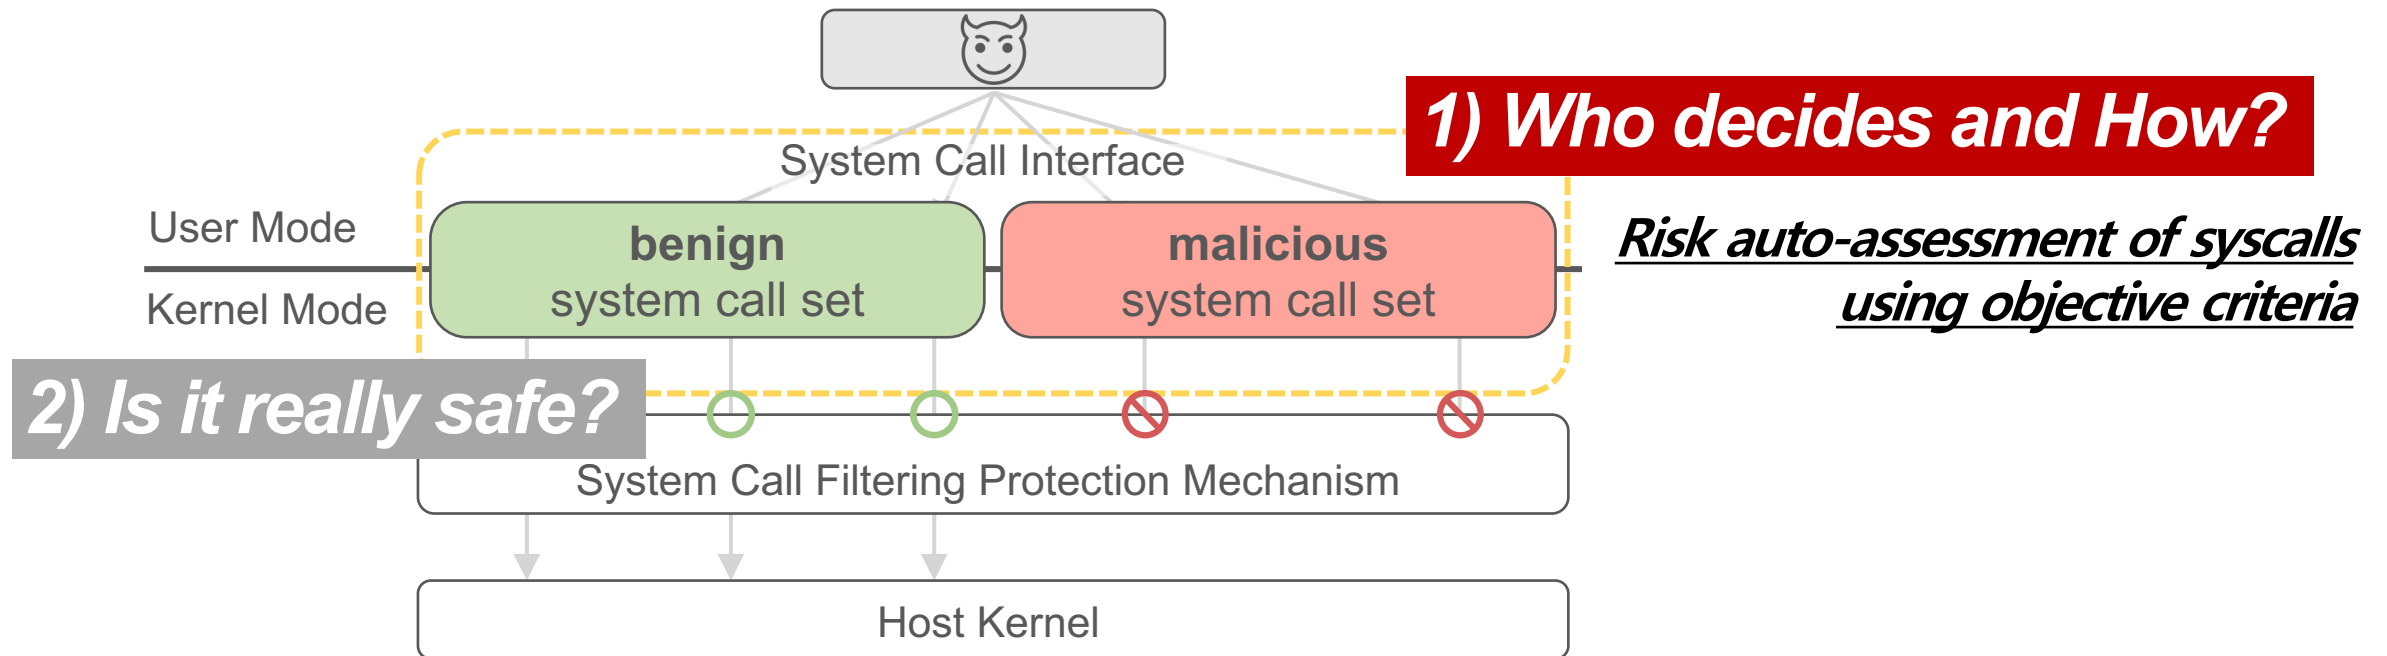
(b) ftrace mechanism

- Tracing tool in Linux kernel
- Possible to write directly to log file



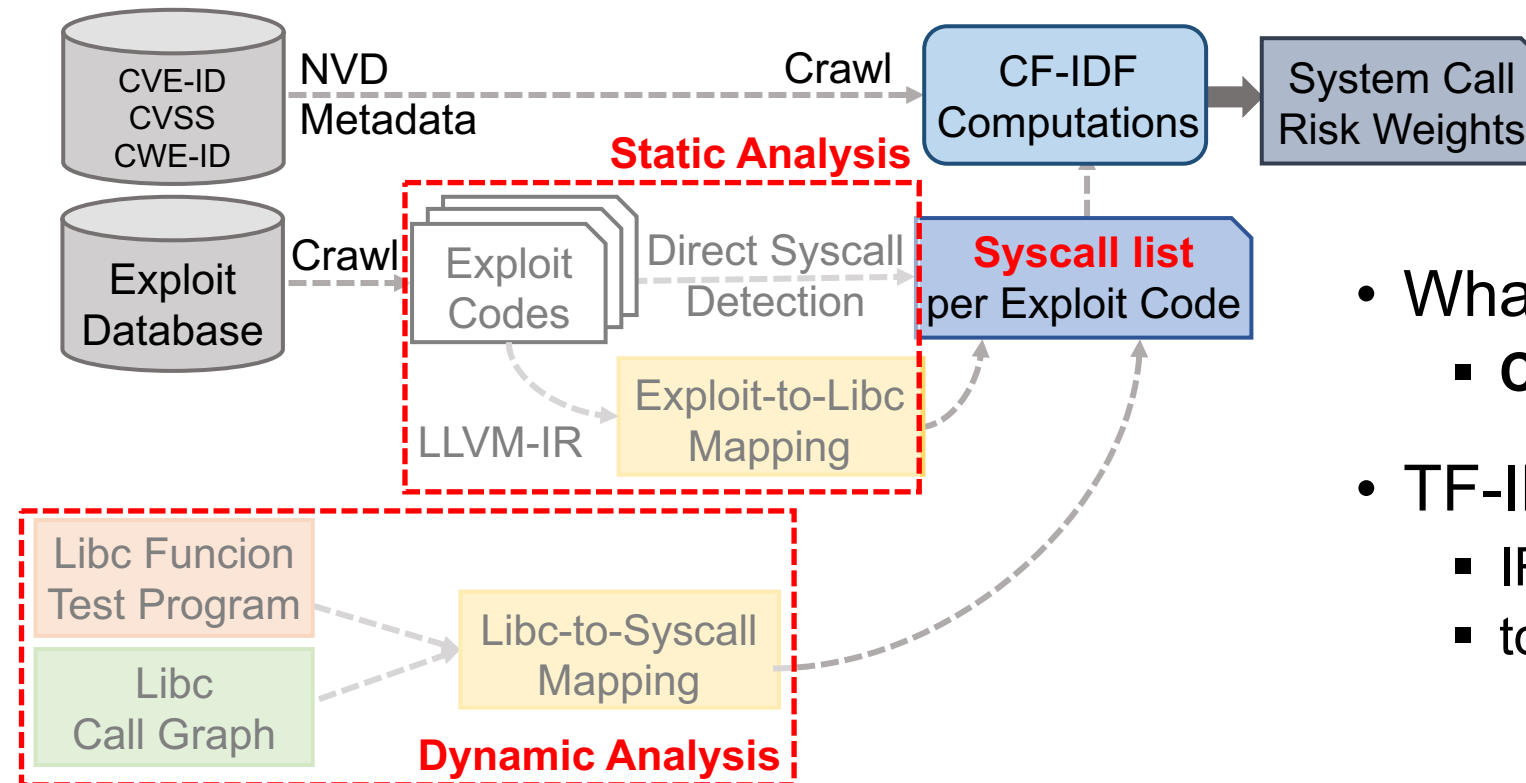
Research on Problem 1

- “Which System Call should be allowed/blocked?”
 - Which System Call is benign/malicious?



Overview of Research on System Call Auto-Assessment of Risk

- Goal: System call risk quantification



- What is a “***Risky System Call***” ?

- Core role in attack logic

- TF-IDF variant = “**CF-IDF**”

- IR technique
- to extract keywords from documents

||
key syscalls

||
exploit codes

Step 1- Input data

Step 2- Exploit Code Analysis

Step 3- System Call Risk
Weight Assignment

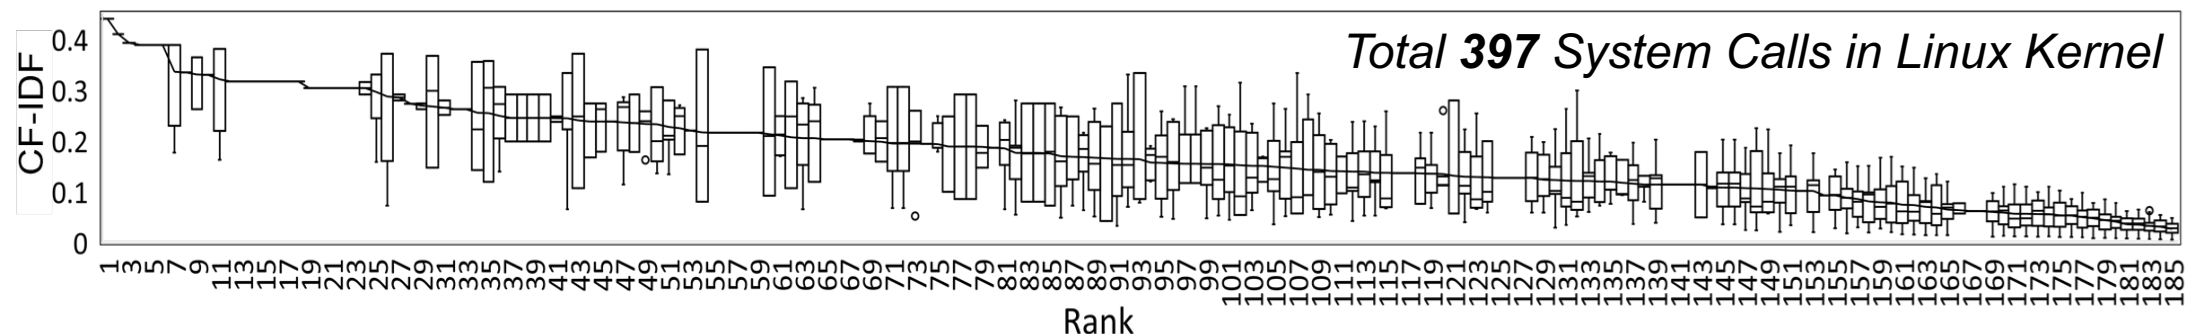
System Call Risk Assessment Result

- The **risk of 185 system** calls was automatically quantified through objective metrics from 298 exploits
- Possible to **filter reflecting the latest vulnerability status**

CVSS of Vulnerability is **10.0** (*highest vulnerability risk score*)

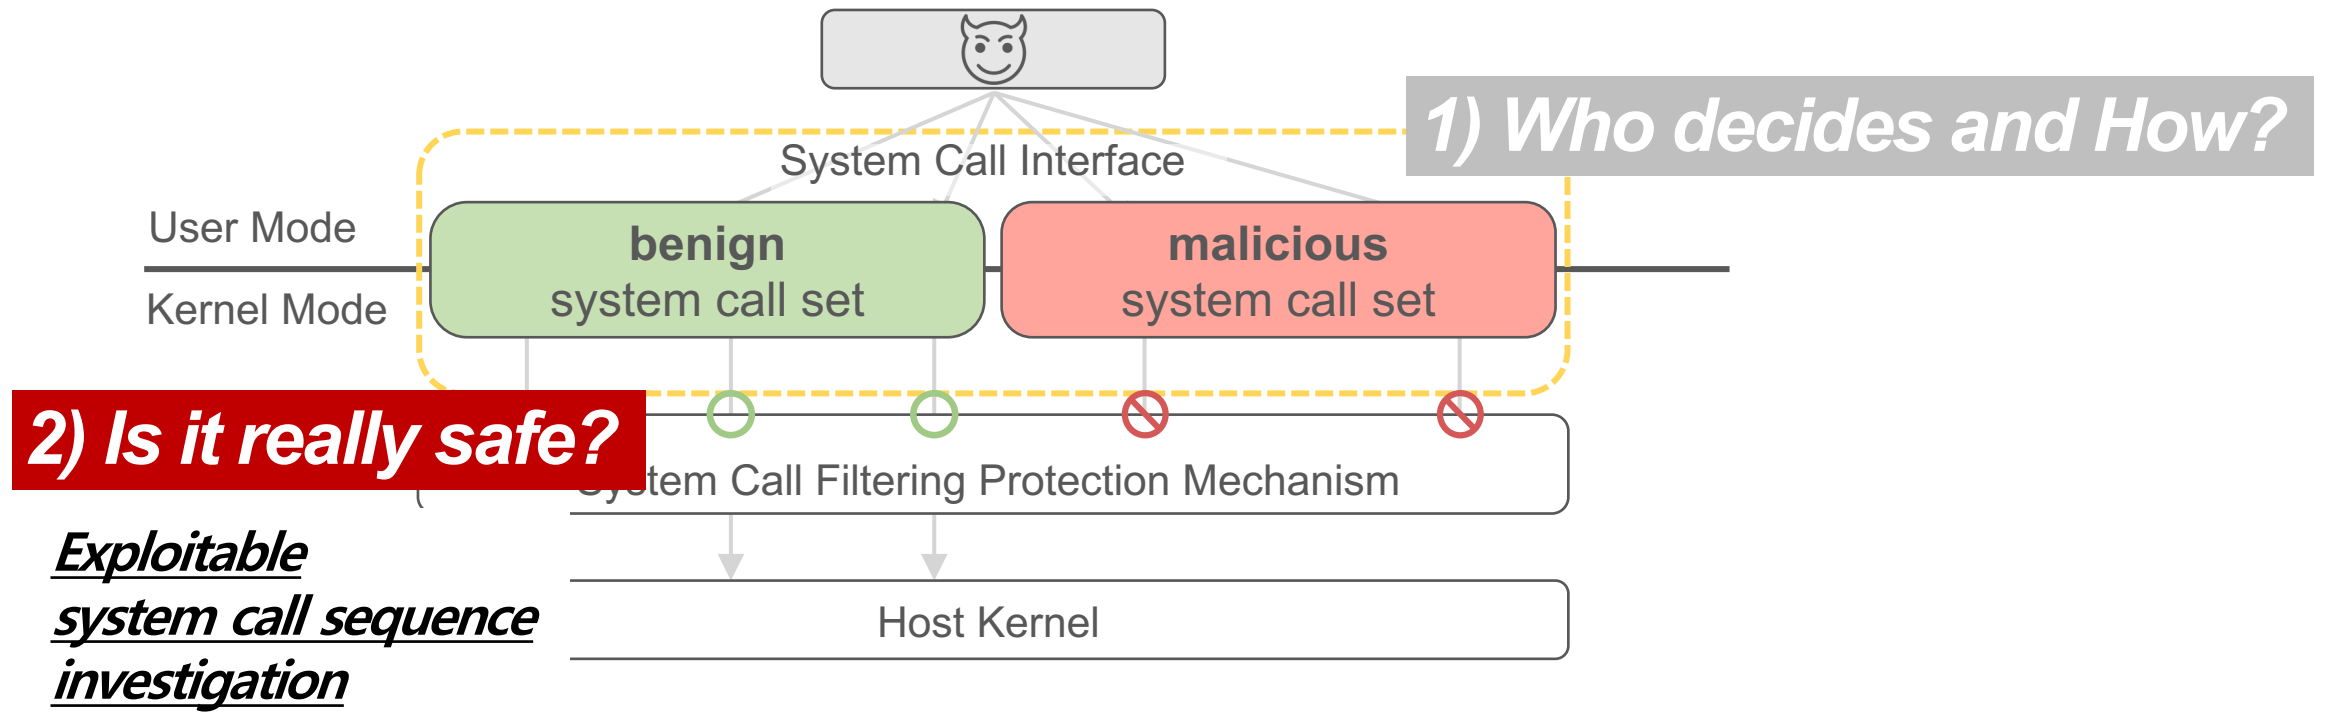
| Rank | Syscall | Weight | Rank | Syscall | Weight | Rank | Syscall | Weight | Rank | Syscall | Weight |
|----------|---------------|-----------------|------|--------------|----------|------|-----------|----------|------------|--------------|-----------------|
| 1 | capset | 0.439551 | 4 | io_uring_reg | 0.388023 | 11 | keyctl | 0.32028 | 12 | io_setup | 0.316477 |
| 2 | add_key | 0.4094431 | 7 | shutdown | 0.335366 | 12 | fchown | 0.316477 | 12 | io_submit | 0.316477 |
| 3 | recvmmsg | 0.392371 | 8 | settimeofday | 0.334059 | 12 | flock | 0.316477 | 12 | kcmp | 0.316477 |
| 4 | getresuid | 0.388023 | 9 | rename | 0.329819 | 12 | mknod | 0.316477 | ... | ... | ... |
| 4 | sendfile | 0.388023 | 10 | creat | 0.329663 | 12 | mq_notify | 0.316477 | 185 | close | 0.026151 |

Appears in **192** out of 298 exploits (about **64%**)



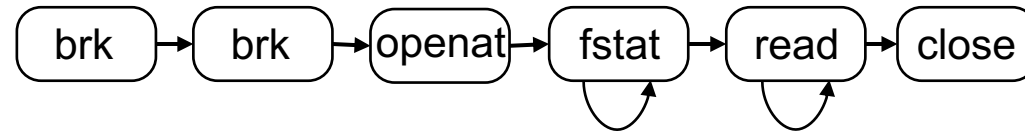
Research on Problem 2

- “Which System Call should be allowed/blocked?”
 - Which System Call is benign/malicious?



Motivation of Malicious System Call Sequence Research

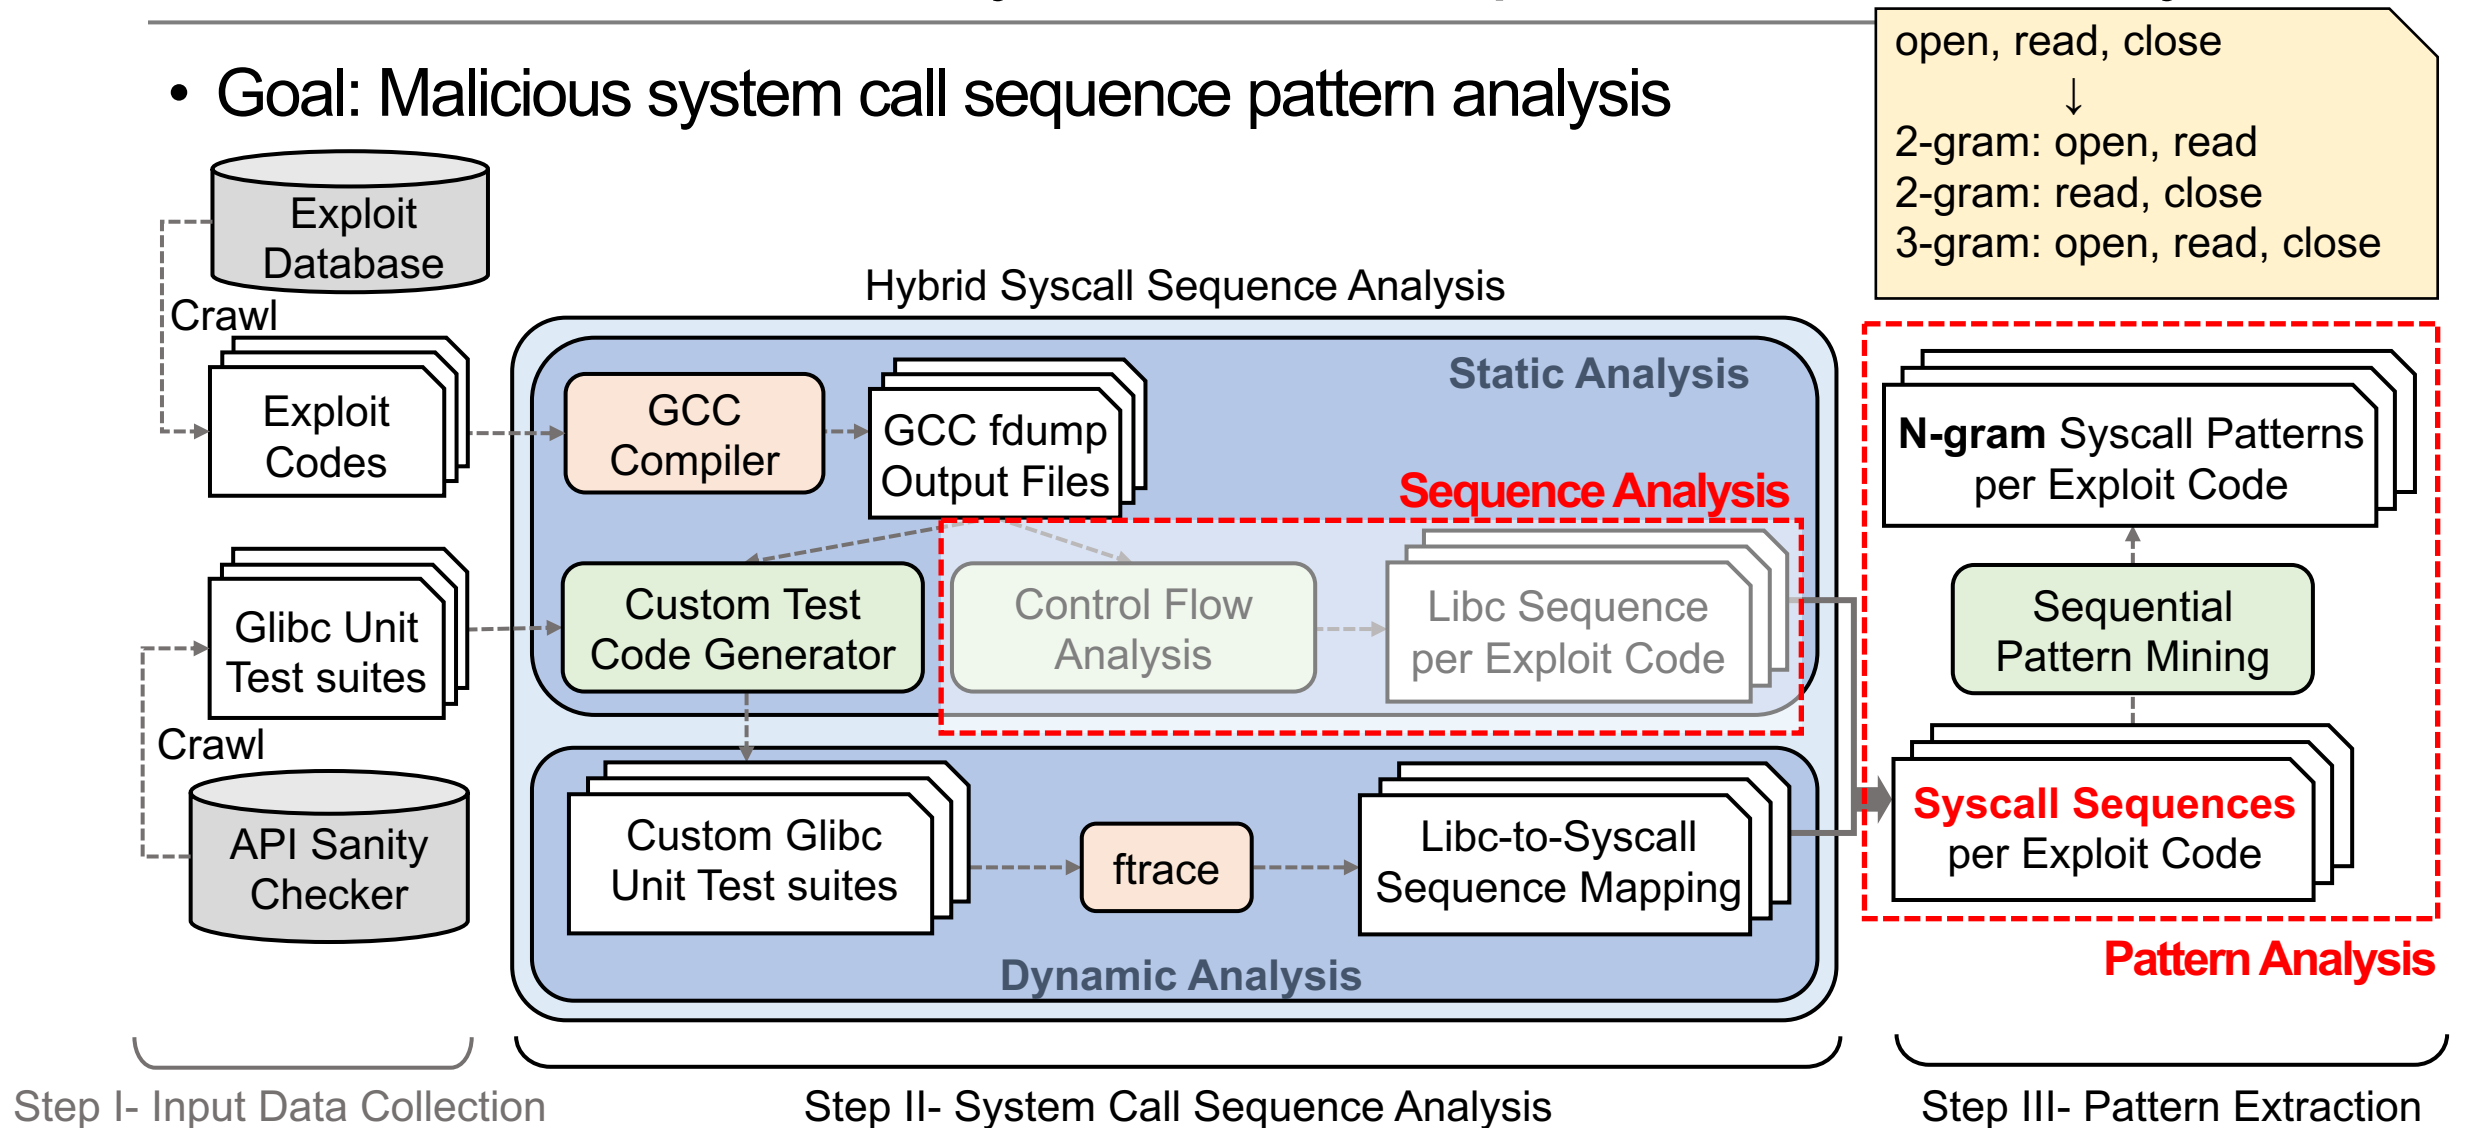
- A malicious sequence consisting of only allowed system calls



- Motivation: Limitation of Docker's Seccomp Profile
- Docker **CAN NOT block 55%** of 106 exploits
 - The 47 host kernel vulnerabilities CAN be exploited using only system calls allowed by Docker's Seccomp profile
- The proposed mechanism **CAN block 70% of exploits that Docker cannot block**
- Towards **system call sequence filtering** for enhanced container security
 - "Can we find patterns in the malicious system call sequences?"

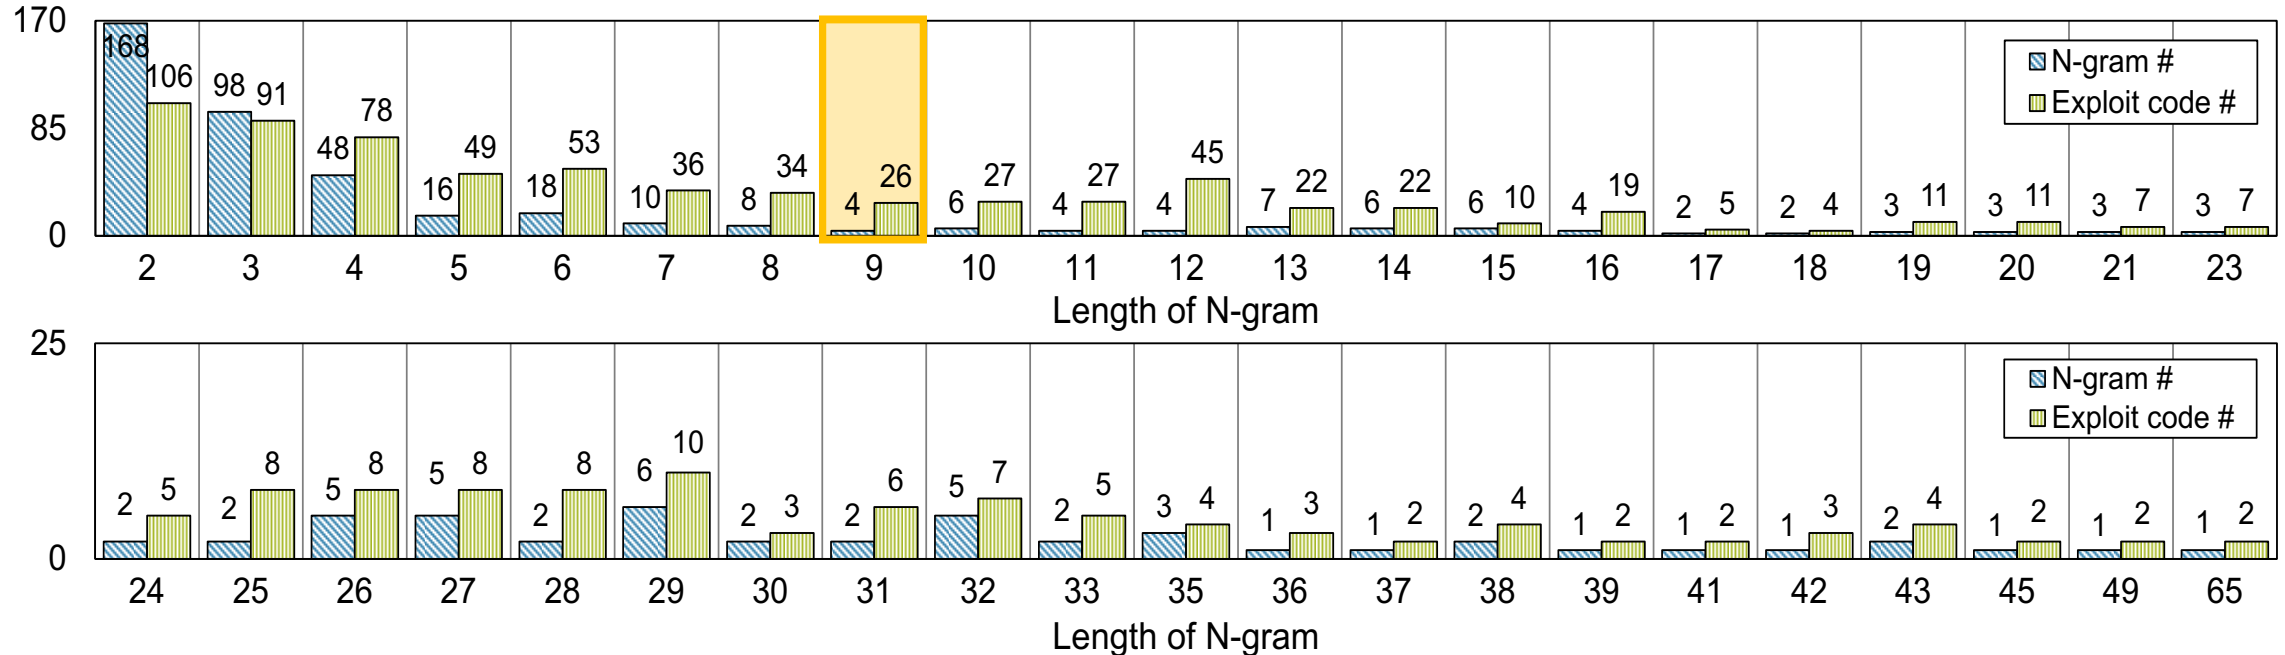
Overview of Malicious System Call Sequence Pattern Analysis

- Goal: Malicious system call sequence pattern analysis



Malicious System Call Sequence Pattern Result

- The 471 system call sequence patterns are found from 106 exploits



| 9-gram # | 9-gram pattern | Exploit # | Total Exploits |
|----------|---|-----------|----------------|
| 4 | (openat,stat,clock_nanosleep) ³ | 3 | 26 |
| | (brk,brk,openat) ² ,fstat,read,fstat | 7 | |
| | (fstat,read) ⁴ ,close | 18 | |
| | brk,brk,openat,(fstat,read) ³ | 22 | |

Effectiveness of Sequence-based Filtering Mechanism

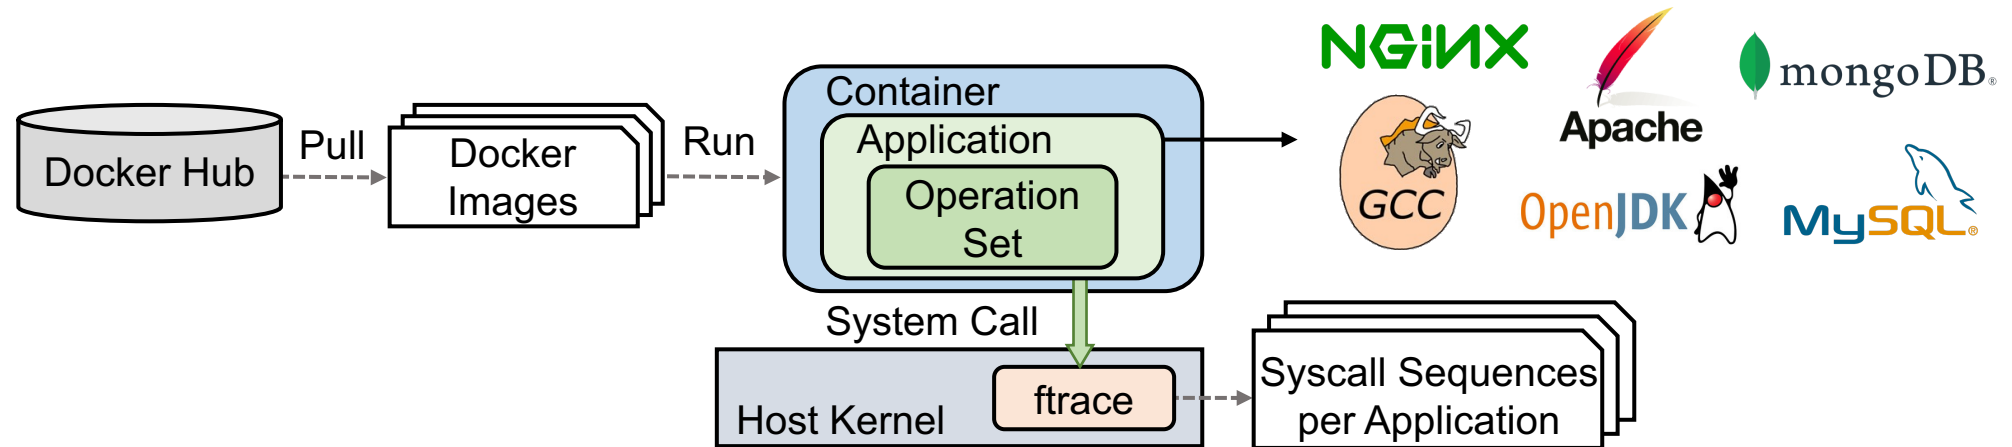
- *What's the effect of pattern lengths on identifying malicious system call sequences?*

| Rank | Length | N-gram pattern | Exploit percentage |
|------|--------|------------------------------|--------------------|
| 1 | 2 | brk, brk | 52% |
| 2 | 2 | openat, fstat | 38% |
| 3 | 2 | read, close | 36% |
| 4 | 2 | rt_sigaction, rt_sigprocmask | 35% |
| 4 | 2 | read, fstat | 33% |

- Top patterns that appear in the most number of exploits
 - provide high coverage when employed as filtering candidates against exploits
 - likely interrupt benign application operations
- Condition for better defense ability of system call sequence pattern
 - **Appearing in multiple exploits**
 - **Longer**

Investigation of Patterns Falsely Disrupting Benign Application

- To what degree do patterns falsely disrupt benign application executions?*



- We have collected per-thread system call traces from **15 popular applications**
 - ▶ Performing diverse application-specific operation
 - e.g., {GET, PUT, POST} for web servers
 - e.g., {INSERT, DELETE, SELECT} for NoSQL databases, amongst others

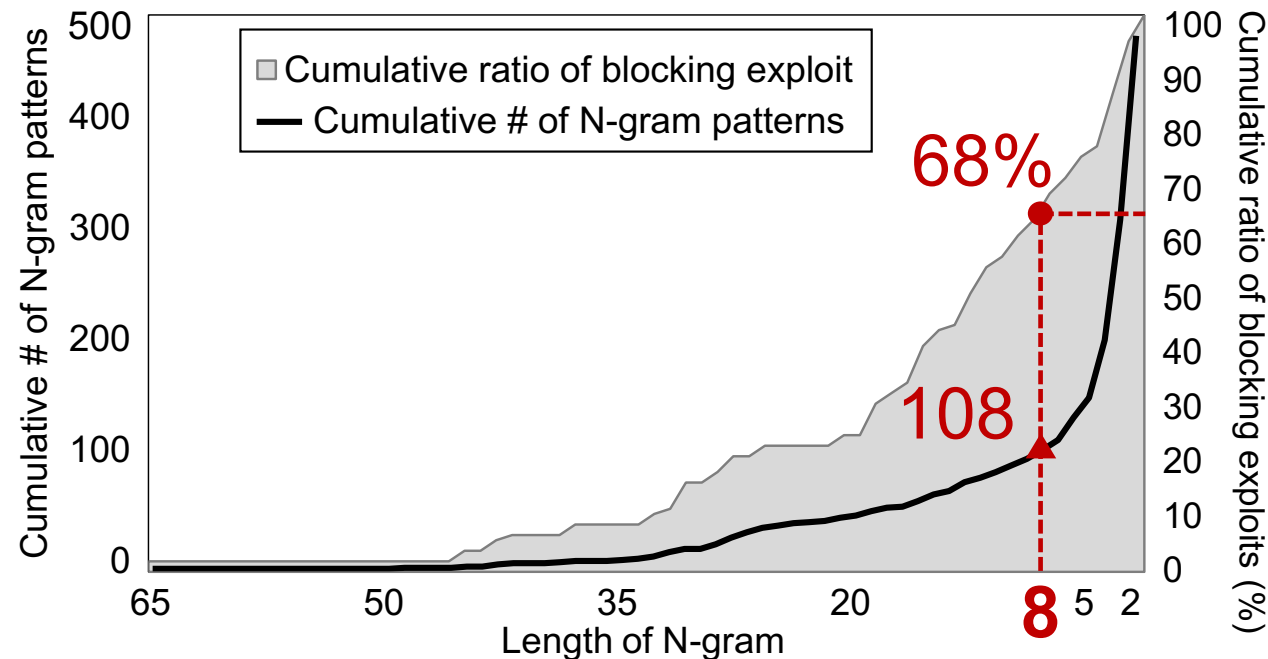
Comparison of Malicious Patterns and Normal Applications

- There exist certain length of N-grams **on the low side** above which do not disrupt normal applications

| Application Name | # of N-grams from exploit codes found in application trace (%) | | | | | | | | | |
|-----------------------------------|--|------------|-----------|-----------|-----------|---------|--------|-----|---------|-------------|
| | 2-gram | 3-gram | 4-gram | 5-gram | 6-gram | 7-gram | 8-gram | ... | 65-gram | Total |
| nginx | 3 (1.8%) | 1 (1%) | – | – | – | – | – | ... | – | 4 (0.8%) |
| httpd | 12 (7.1%) | 2 (2%) | – | – | – | – | – | ... | – | 14 (3%) |
| tomcat | 7 (4.2%) | 3 (3.1%) | – | – | – | – | – | ... | – | 10 (2.1%) |
| node | 4 (2.4%) | 2 (2%) | – | – | – | – | – | ... | – | 6 (1.3%) |
| mongodb | 19 (11.3%) | 6 (6.1%) | 1 (2.1%) | – | – | – | – | ... | – | 26 (5.5%) |
| mysql | 22 (13.1%) | 11 (11.2%) | 5 (10.4%) | 2 (12.5%) | – | – | – | ... | – | 40 (8.5%) |
| mariadb | 11 (6.5%) | 2 (2%) | – | – | – | – | – | ... | – | 13 (2.8%) |
| redis | 9 (5.4%) | 3 (3.1%) | – | – | 1 (5.6%) | 1 (10%) | – | ... | – | 14 (3%) |
| gcc | 34 (20.2%) | 16 (16.3%) | 5 (10.4%) | 3 (18.8%) | 1 (5.6%) | – | – | ... | – | 59 (12.5%) |
| openjdk | 50 (29.8%) | 18 (18.4%) | 5 (10.4%) | 3 (18.8%) | 2 (11.1%) | – | – | ... | – | 78 (16.6%) |
| gzip | 9 (5.4%) | 13 (13.3%) | 4 (8.3%) | – | 1 (5.6%) | – | – | ... | – | 27 (5.7%) |
| bzip2 | 9 (5.4%) | 8 (8.2%) | 3 (6.3%) | 1 (6.3%) | 1 (5.6%) | – | – | ... | – | 22 (4.7%) |
| qalc | 48 (28.6%) | 14 (14.3%) | 5 (10.4%) | 3 (18.8%) | 2 (11.1%) | – | – | ... | – | 72 (15.3%) |
| ghostscript | 36 (21.4%) | 21 (21.4%) | 6 (12.5%) | 2 (12.5%) | 1 (5.6%) | – | – | ... | – | 66 (14%) |
| lowriter | 71 (42.3%) | 28 (28.6%) | 7 (14.6%) | 4 (25%) | 2 (11.1%) | – | – | ... | – | 112 (23.8%) |
| Total | 82 (48.8%) | 45 (45.9%) | 9 (18.8%) | 4 (25%) | 3 (16.7%) | 1 (10%) | – | ... | – | 144 (30.6%) |
| # of N-grams from exploits | 168 | 98 | 48 | 16 | 18 | 10 | 8 | ... | 1 | 471 |
| Proportion of blocked apps | 100% | 100% | 60% | 46.6% | 53.3% | 6% | – | – | – | – |

Short Sequence is Good Enough!

- Blocking capability and the ratio that blocks benign applications are in trade-off relationship
 - From the application set we experimented with, the system call sequence length of 8 and above has been shown to be **free from disruptive side-effects**
 - The use of all the N-grams starting from **length 8** and above up to 65-gram **block 68%** of 108 exploit codes we tested



All of attacks that Docker Seccomp CANNOT prevent can block!

- Docker Seccomp **CAN NOT** protect against **58 (55%)** out of the 106 exploits
- Patterns with a length of 8 or more **can mitigate 40** out of the 58 exploits (about **70% Coverage**)
- For the remaining 18 (30%) exploits that share only short patterns (lengths 1 through 7), it is still possible to mitigate them by using exploit-specific system call sequences that are not caught as patterns
 - Candidate with length of 13 for CVE-2022-0847 exploit:
< openat, fstat, pipe, fcntl, write³, read³, splice, write, close >

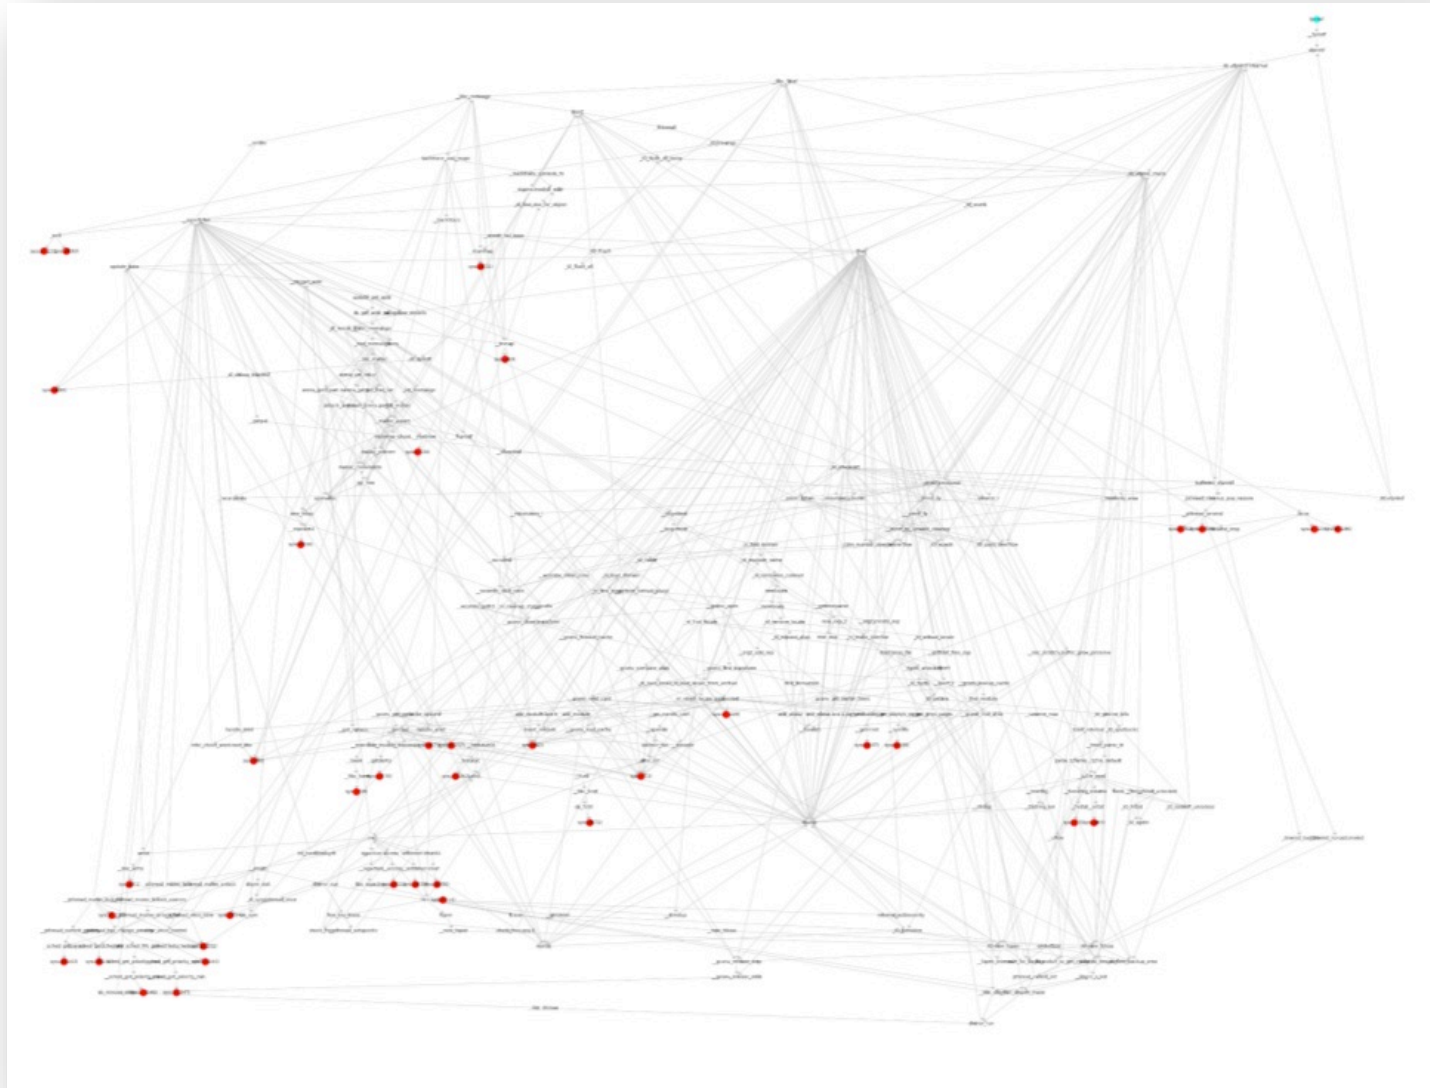
Conclusion

- System call filtering protection is important for protection of container environment
 - Attack surface of shared kernel reduction
 - Huge damage caused by container escape
- The current filtering mechanism is fragile and non-scalable solution
- The challenge to select an appropriate system calls to block
- System call risk reflecting the latest vulnerability status can be a guide for system call policies
- System call sequence blocking can compensate for the loopholes in individual system call blocking
- *Prevention of software disasters is essential as well as recovery*

Reference

1. <https://financesonline.com/cloud-computing-statistics/>.
2. "Seccomp profiles for Docker," <https://docs.docker.com/engine/security/seccomp/>.
3. Seyedhamed Ghavamnia, Tapti Palit, Azzedine Benameur, and Michalis Polychronakis. 2020. Confine: Automated System Call Policy Generation for Container Attack Surface Reduction. In 23rd International Symposium on Research in Attacks, Intrusions and Defenses (RAID 2020).
4. Nicholas DeMarinis, Kent Williams-King, Di Jin, Rodrigo Fonseca, and Vasileios P. Kemerlis. 2020. sysfilter: Automated System Call Filtering for Commodity Software. In 23rd International Symposium on Research in Attacks, Intrusions and Defenses (RAID 2020).
5. Muhammad Abubakar, Adil Ahmad, Pedro Fonseca, and Dongyan Xu. 2021. SHARD: Fine-Grained Kernel Specialization with Context-Aware Hardening. In USENIX Security Symposium.
6. Olufogorehan Tunde-Onadele, Yuhang Lin, Jingzhu He, and Xiaohui Gu. 2020. Self-Patch: Beyond Patch Tuesday for Containerized Applications. In IEEE ACSOS.
7. Nuno Lopes, Rolando Martins, Manuel Eduardo Correia, Sérgio Serrano, and Francisco Nunes. 2020. Container Hardening Through Automated Seccomp Profiling. In Proceedings of the 2020 6th International Workshop on Container Technologies and Container Clouds.
8. Wan, Zhiyuan, et al. "Mining sandboxes for linux containers." 2017 IEEE International Conference on Software Testing, Verification and Validation (ICST).
9. L. Lei et al., "SPEAKER: Split-Phase Execution of Application Containers," in Detection of Intrusions and Malware, and Vulnerability Assessment.
10. API Sanity Checker. An automatic generator of basic unit tests for a shared C/C++ library. <https://lvc.github.io/api-sanity-checker/>.

Appendix. fprintf glibc library function call graph



in Confine [3]

Static Analysis to Map Exploit Code-Glibc Function

- Using Clang + LLVM-IR
 - Getting info on invocation of glibc function and direct system call

```
inline _syscall1(int, close, int, fd); ① direct syscall (inline)

void alloc_victim(void) ② user-defined function
{
    fd = open(FILE_PATH, O_WRONLY, mode); ③ library function
    ④ close(fd); invocation of ①
    __asm__(
        "pusha\n"
        "movl %1, %%eax\n"
        "movl $(\"xstr(CLONEFL)\", %%ebx\n"
        "movl %%esp, %%ecx\n"
        "movl $120, %%eax\n" ⑤ direct syscall (asm)
        "int $0x80\n"
        "movl %%eax, %0\n"
        "popa\n"
        : : "m" (pid), "m" (dummy)
    );
    syscall(SYS_exit, 0); ⑥ direct syscall (with syscall())
}
```

Example source code

```
define dso_local void @alloc_victim() ② {
    %1 = load i32, i32* @mode, align 4
    %2 = call i32 @i8*, i32, ... @open(i8* ③
        getelementptr inbounds ([14 x i8], [14 x i8]* @.str, i64
        0, i64 0), i32 1, i32 %1) #6
    ...
    %8 = call i32 @close(i32 %7) ④
    call void asm sideeffect "pusha\0Amovl $1, %eax\0Amovl
    $(CLONEFL), %ebx\0Amovl %esp, %ecx\0A
    movl $$120, %eax\0Aint $$0x80\0Amovl %eax, $0\0Apopa\0A",
    "m,m,~{dirflag},~{fpsr},~{flags}"(i32* @pid, i32*
    @dummy) #4, !srcloc !3
    %9 = call i64 (i64, ...) @syscall(i64 60, i32 0) #7 ⑥
    br label %10
10:
    ret void
}

define dso_local i32 @close(i32 %0) #0 {
    ...
    %6 = call i64 asm sideeffect "syscall", "={ax},0,{di},
    ~{r11},~{rcx},~{memory},~{dirflag},~{fpsr},~{flags}" ①
    (i64 3, i64 %5) #4, !srcloc !2
    ...
}
```

LLVM-IR format

Appendix. CF-IDF Dtails

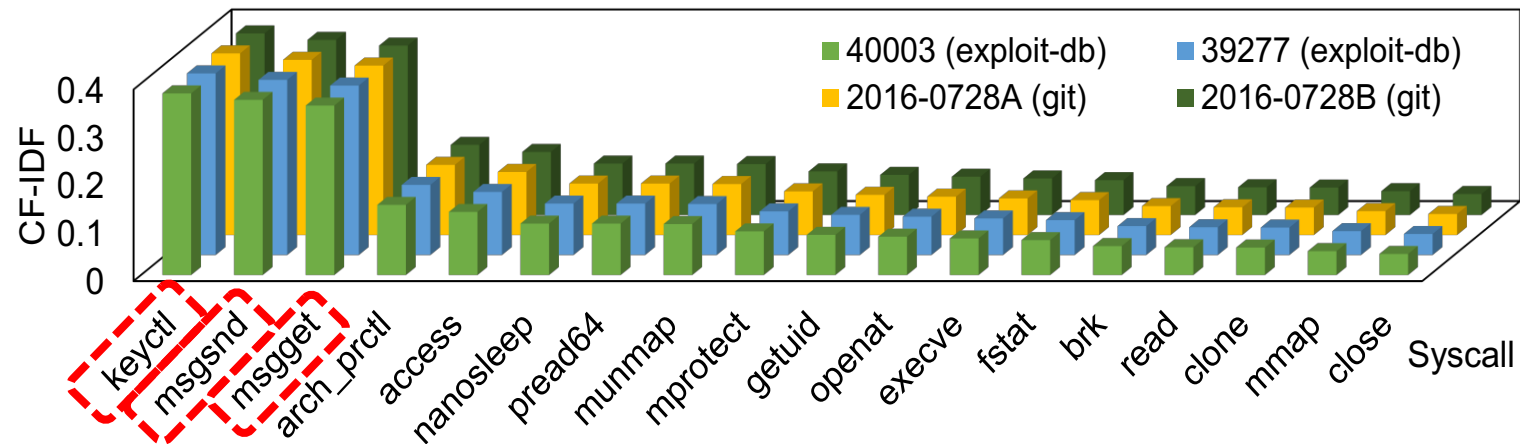
- TF-IDF variant
 - Document = Exploit Code
 - Term = System Call
 - **IDF(Inverse Document Frequency)**
$$= \frac{|Doc Set|}{Term f in Doc Set}$$
 - ▶ TRUE: close, exit, nanosleep in most codes
 - TF(Term Frequency)
$$= Term f in a Doc$$
 - ▶ FALSE
 - **CF(Class Frequency)**
$$= Doc f with Term in Class$$
 - ▶ Class = Vulnerability
 - ▶ pipe & splice always in Dirty Pipe
- $CF-IDF = CF \times IDF$

Appendix. Full Ranked List of System Calls

| Rank | System call | Weight | Rank | System call | Weight | Rank | System call | Weight | Rank | System call | Weight | Rank | System call | Weight |
|------|--------------------|----------|------|-------------------|----------|------|-------------------|----------|------|-----------------|----------|------|----------------|----------|
| 1 | capset | 0.439551 | 37 | shmat | 0.244536 | 75 | mkdir | 0.193169 | 112 | pipe | 0.139035 | 149 | setgid | 0.104811 |
| 2 | add_key | 0.409431 | 37 | sigaltstack | 0.244536 | 76 | getdents64 | 0.188096 | 113 | getppid | 0.138948 | 150 | connect | 0.103625 |
| 3 | recvmmsg | 0.392371 | 37 | setxattr | 0.244536 | 77 | llseek | 0.187949 | 114 | sendmsg | 0.136639 | 151 | prlimit64 | 0.101779 |
| 4 | getresuid | 0.388023 | 41 | symlink | 0.244057 | 77 | getpriority | 0.187949 | 115 | setresgid | 0.136633 | 152 | seccomp | 0.101097 |
| 4 | sendfile | 0.388023 | 42 | getcwd | 0.244007 | 79 | setns | 0.187368 | 116 | uselib | 0.136261 | 153 | stat | 0.100942 |
| 4 | io_uring_register | 0.388023 | 43 | fchmod | 0.240128 | 80 | msgrcv | 0.185928 | 116 | msync | 0.136261 | 154 | setitimer | 0.092306 |
| 7 | shutdown | 0.335366 | 44 | modify_ldt | 0.237571 | 81 | getsockname | 0.184515 | 118 | mincore | 0.135555 | 155 | setsockopt | 0.091907 |
| 8 | settimeofday | 0.334059 | 44 | clock_gettime | 0.237571 | 82 | setrlimit | 0.175967 | 119 | unname | 0.135125 | 156 | lseek | 0.087384 |
| 9 | rename | 0.329819 | 46 | process_vm_readv | 0.237358 | 83 | getrlimit | 0.175651 | 120 | pause | 0.134133 | 157 | wait4 | 0.084849 |
| 10 | creat | 0.329663 | 47 | writev | 0.235578 | 83 | sync | 0.175651 | 121 | vmsplice | 0.130688 | 158 | exit_group | 0.080511 |
| 11 | keyctl | 0.32028 | 48 | getdents | 0.234431 | 85 | splice | 0.174387 | 122 | alarm | 0.128849 | 159 | getpid | 0.078393 |
| 12 | fchown | 0.316477 | 49 | sendmmsg | 0.232625 | 86 | ptrace | 0.169211 | 123 | setresuid | 0.128281 | 160 | ioctl | 0.077965 |
| 12 | flock | 0.316477 | 50 | syslog | 0.232287 | 87 | setpriority | 0.167943 | 124 | gettid | 0.128004 | 161 | arch_prctl | 0.073672 |
| 12 | mknod | 0.316477 | 51 | mount | 0.226888 | 88 | userfaultfd | 0.167438 | 125 | epoll_create1 | 0.126591 | 162 | rt_sigaction | 0.072893 |
| 12 | mq_notify | 0.316477 | 52 | rmdir | 0.224417 | 89 | futex | 0.166319 | 125 | setgroups | 0.126591 | 163 | kill | 0.069353 |
| 12 | io_setup | 0.316477 | 53 | getgroups | 0.219776 | 90 | statfs | 0.165307 | 125 | umask | 0.126591 | 164 | access | 0.068628 |
| 12 | io_submit | 0.316477 | 54 | select | 0.216088 | 91 | dup2 | 0.164186 | 128 | unlink | 0.126579 | 165 | exit | 0.06426 |
| 12 | kcmp | 0.316477 | 55 | pwrite64 | 0.21538 | 92 | accept | 0.164048 | 129 | time | 0.123122 | 166 | renameat2 | 0.063323 |
| 19 | futimesat | 0.30329 | 55 | set_mempolicy | 0.21538 | 93 | perf_event_open | 0.163461 | 130 | socketpair | 0.122321 | 167 | sysinfo | 0.061537 |
| 19 | inotify_rm_watch | 0.30329 | 55 | readv | 0.21538 | 94 | poll | 0.15674 | 131 | geteuid | 0.121388 | 167 | setreuid | 0.061537 |
| 19 | inotify_init1 | 0.30329 | 55 | sched_getaffinity | 0.21538 | 95 | getsockopt | 0.15602 | 132 | setuid | 0.120857 | 169 | socket | 0.059526 |
| 19 | restart_syscall | 0.30329 | 55 | shmdt | 0.21538 | 96 | sched_setaffinity | 0.155237 | 133 | fruncate | 0.1204 | 170 | rt_sigprocmask | 0.058141 |
| 19 | utimensat | 0.30329 | 60 | mremap | 0.212287 | 97 | timerfd_create | 0.154563 | 134 | mlock | 0.119328 | 171 | pread64 | 0.055834 |
| 24 | clock_nanosleep | 0.303143 | 61 | inotify_init | 0.210391 | 97 | timerfd_settime | 0.154563 | 135 | setsid | 0.119298 | 172 | munmap | 0.055819 |
| 25 | umount2 | 0.295618 | 62 | sched_yield | 0.206736 | 99 | unshare | 0.153808 | 136 | epoll_ctl | 0.116779 | 173 | write | 0.055482 |
| 26 | chown | 0.286681 | 63 | recvmmsg | 0.205074 | 100 | fcntl | 0.153628 | 137 | sendto | 0.115332 | 174 | nanosleep | 0.055054 |
| 27 | link | 0.284955 | 64 | getegid | 0.204785 | 101 | madvise | 0.152935 | 138 | setpgid | 0.113589 | 175 | open | 0.05293 |
| 28 | dup3 | 0.272522 | 65 | fallocate | 0.202194 | 102 | gettimeofday | 0.151181 | 139 | getgid | 0.113317 | 176 | getuid | 0.052282 |
| 29 | eventfd2 | 0.269163 | 65 | sysctl | 0.202194 | 103 | tgkill | 0.150511 | 140 | getresgid | 0.113173 | 177 | mprotect | 0.049306 |
| 30 | msgsnd | 0.267191 | 65 | move_pages | 0.202194 | 104 | personality | 0.150292 | 140 | adjtimex | 0.113173 | 178 | execve | 0.04722 |
| 31 | sched_setscheduler | 0.264745 | 68 | shmctl | 0.200075 | 105 | listen | 0.148191 | 140 | timer_create | 0.113173 | 179 | openat | 0.043193 |
| 32 | inotify_add_watch | 0.261581 | 69 | msgctl | 0.199029 | 106 | prctl | 0.146587 | 140 | memfd_create | 0.113173 | 180 | fstat | 0.040249 |
| 32 | waitid | 0.261581 | 70 | dup | 0.197707 | 107 | readlink | 0.144822 | 144 | epoll_wait | 0.108171 | 181 | clone | 0.035285 |
| 34 | msgget | 0.254518 | 71 | io_uring_enter | 0.194227 | 108 | chroot | 0.142477 | 145 | set tid address | 0.107822 | 182 | brk | 0.034785 |
| 35 | pipe2 | 0.25433 | 71 | io_uring_setup | 0.194227 | 109 | bpf | 0.142335 | 145 | set_robust_list | 0.107822 | 183 | read | 0.032339 |
| 36 | chmod | 0.248783 | 73 | chdir | 0.19396 | 110 | recvfrom | 0.140137 | 147 | bind | 0.106639 | 184 | mmap | 0.03019 |
| 37 | shmget | 0.244536 | 74 | iopl | 0.193403 | 111 | epoll_create | 0.139559 | 148 | rt_sigreturn | 0.1059 | 185 | close | 0.026151 |

Appendix. Verification of Our Risk Metric

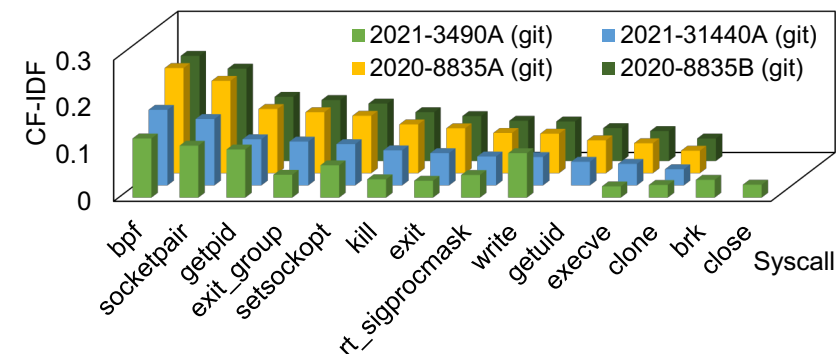
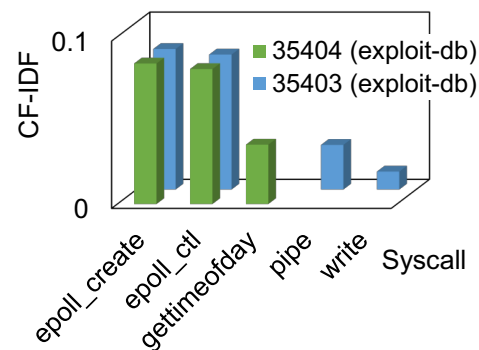
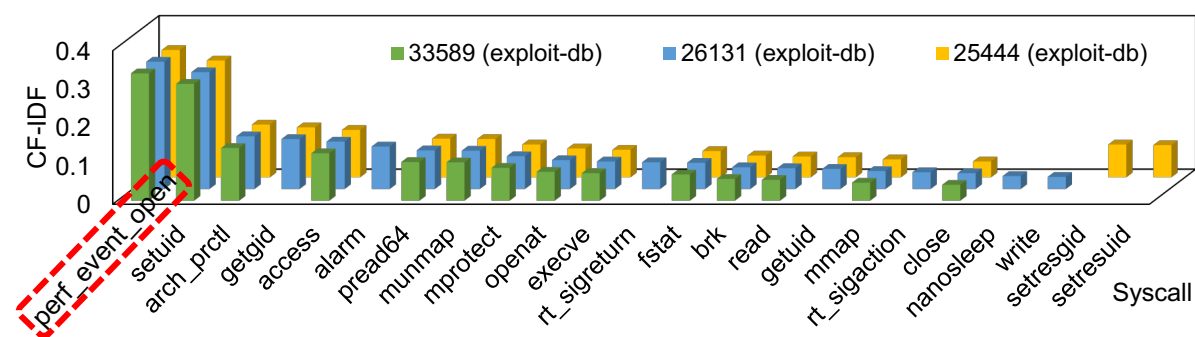
- Empirical justification supporting the soundness of the CF-IDF metric
 - No previous work that suggested the quantified risk of system calls
- CF-IDF risk vectors are signatures that represent the characteristic of exploit codes
 - It would generate similar signatures for the exploit codes that are truly similar in their nature



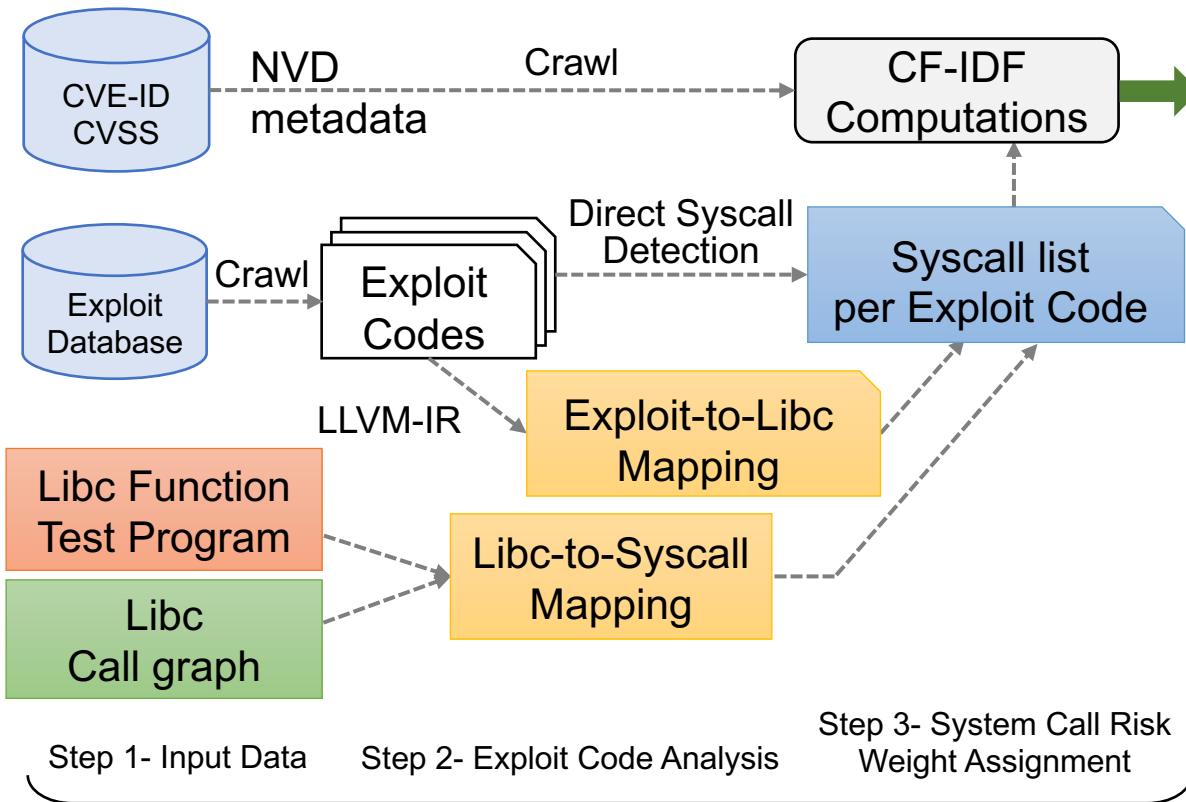
| Exploit-ID (source) | CVE-ID | CVE Description | Cosine Similarity |
|---------------------|---------------|---|-------------------|
| 40003 (exploit-db) | CVE-2016-0728 | Keyring object reference mishandling with crafted keyctl syscall | - |
| 39277 (exploit-db) | | | 1.0 |
| 2016-0728A (git) | | | 1.0 |
| 2016-0728B (git) | | | 1.0 |

Appendix. Additional Sample Groups Similarity Comparison

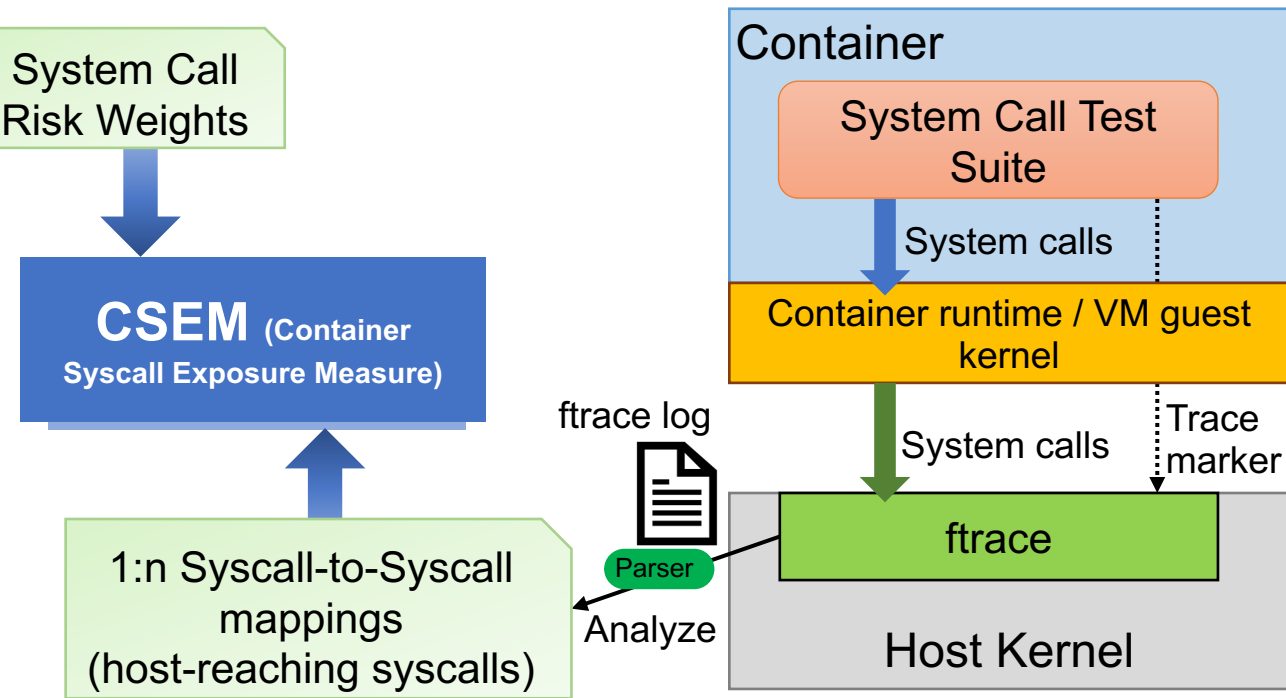
| Group | Grouping Criteria | Exploit-ID (source) | CVE-ID | Similarity |
|-------|---|---------------------|----------------|------------|
| G-I | Same CVE. Incorrect integer data type via crafted perf_event_open. PrivEsc. | 33589 (exploit-db) | CVE-2013-2094 | - |
| | | 26131 (exploit-db) | | 1.0 |
| | | 25444 (exploit-db) | | 1.0 |
| G-II | Different CVE but same vulnerability type. Improper traversal via crafted epoll_create and epoll_ctl. DoS. | 35403 (exploit-db) | CVE-2011-1083 | - |
| | | 35404 (exploit-db) | CVE-2011-1082 | 0.764533 |
| G-III | Different CVE but same vulnerability type. Lack of validation with bpf. PrivEsc. | 2021-31440A (git) | CVE-2021-31440 | - |
| | | 2020-8835A (git) | CVE-2020-8835 | 1.0 |
| | | 2020-8835B (git) | | 1.0 |
| | | 2021-3490A (git) | CVE-2021-3490 | 0.975811 |



Appendix. SecQuant Overview



(a) SCAR (System Call Assessment of Risk)



(b) SCED (System Call Exposure Discovery)

Appendix. Docker Seccomp Profile

- Docker uses the Seccomp kernel feature to block container access to 51 system calls
 - Docker's Seccomp profile[2] allows 346 system calls

| | | | |
|-----------------|-------------------|-------------------|-------------|
| acct | iopl | personality | swapon |
| add_key | kcmp | pivot_root | swapoff |
| bpf | kexec_file_load | process_vm_readv | sysfs |
| clock_adjtime | kexec_load | process_vm_writev | _sysctl |
| clock_settime | keyctl | ptrace | umount |
| clone | lookup_dcookie | query_module | umount2 |
| create_module | mbind | quotactl | unshare |
| delete_module | mount | reboot | uselib |
| finit_module | move_pages | request_key | userfaultfd |
| get_kernel_syms | name_to_handle_at | set_mempolicy | ustat |
| get_mempolicy | nfsservctl | setns | vm86 |
| init_module | open_by_handle_at | settimeofday | vm86old |
| ioperm | perf_event_open | stime | |