SecQuant: Quantifying Container System Call Exposure

Sunwoo Jang[1], **Somin Song**[1], Byungchul Tak[1], Sahil Suneja[2], Michael V. Le[2], Chuan Yue[3], Dan Williams[4]

[1]Kyungpook National University (KNU), Daegu, Republic of Korea

[2]IBM TJ Watson Research Center, Yorktown Heights, NY, USA

[3]Colorado School of Mines, CO, USA

[4]Virginia Tech, Blacksburg, VA, USA

September 27th, 2022

Security for Container Runtimes

- We focus on the ***container escape (break-out)***
  - Containers = namespaced processes

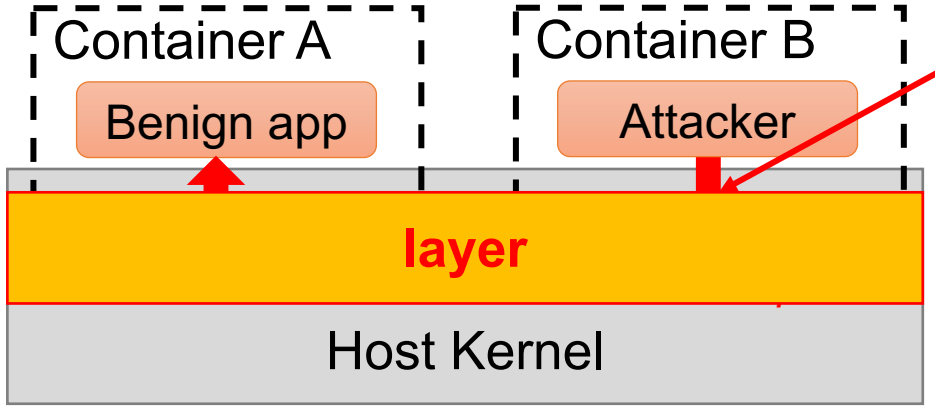

System calls

- Dirty COW (CVE-2016-5195)
  → `write, madvise`
- Dirty Pipe (CVE-2020-0847)
  → `pipe, splice`
- Dirty Cred (CVE-2022-2588/CVE-2021-4154)
  → `writev`

- Handle system calls for the host kernel
  → Smaller attack surface

- Secure Container Runtimes

gVisor (Google) Kata containers Nabla (IBM) RunD [2022 ATC]

Surrogate (Proxy) Layer: gVisor vs. Kata vs. Nabla

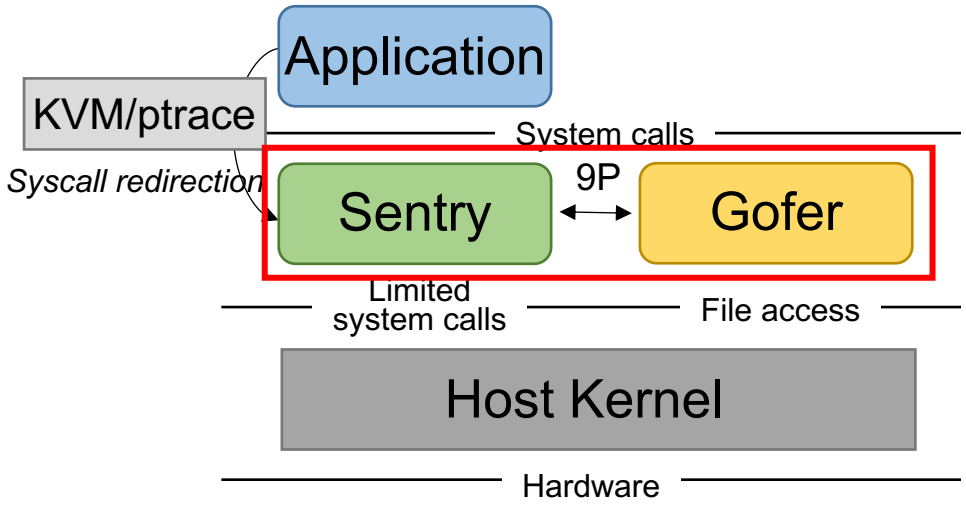

gVisor (Google)

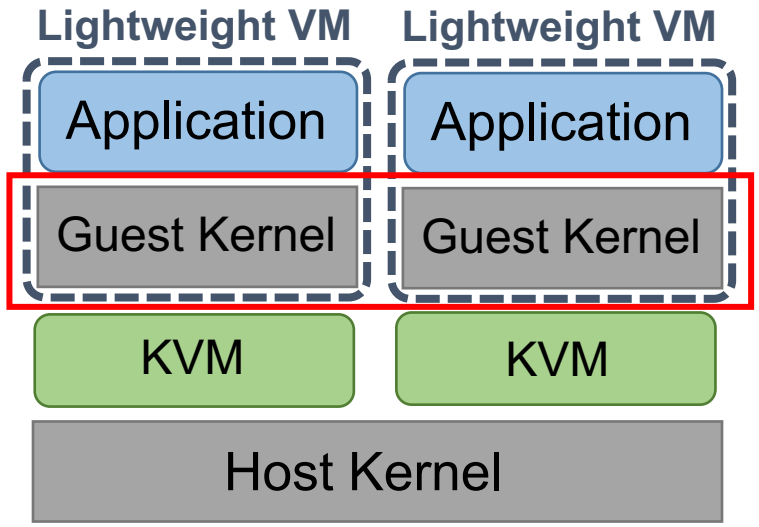

Kata containers

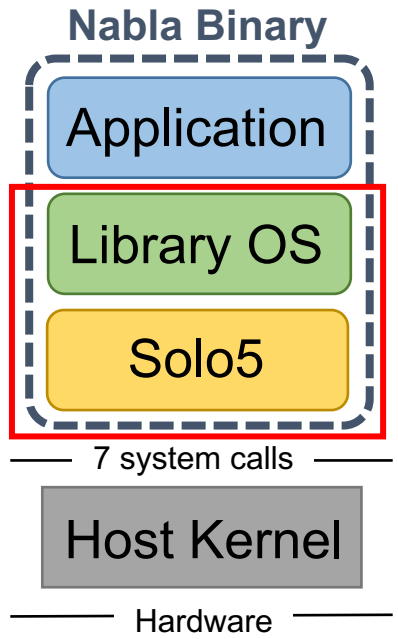

Nabla (IBM)

- **Utility of quantifying container runtime security**
  - Only qualitative statements are given
  - Quantification allows: Comparison, Trend, Engineering, What-if analysis

Our Intuition and Approach

- Attack surface measure: ***system calls***
  - *How many system calls reach the host kernel?*

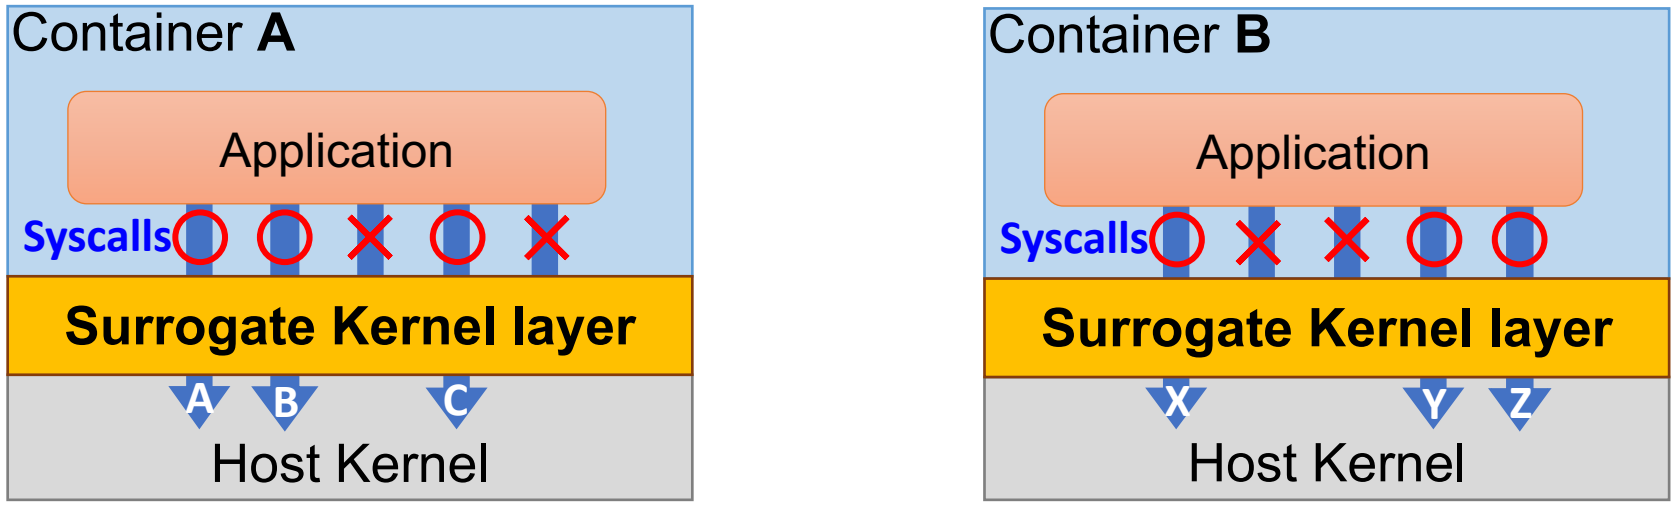

- Simple counting of system calls (and/or types) is insufficient
- Need to determine the **importance of individual system calls**
  - Exploit codes → which system calls are use in the attacks?

Approach Overview

- An unified metric for comparing container security
  - Two perspectives: **Risk** and **Reachability** of system call

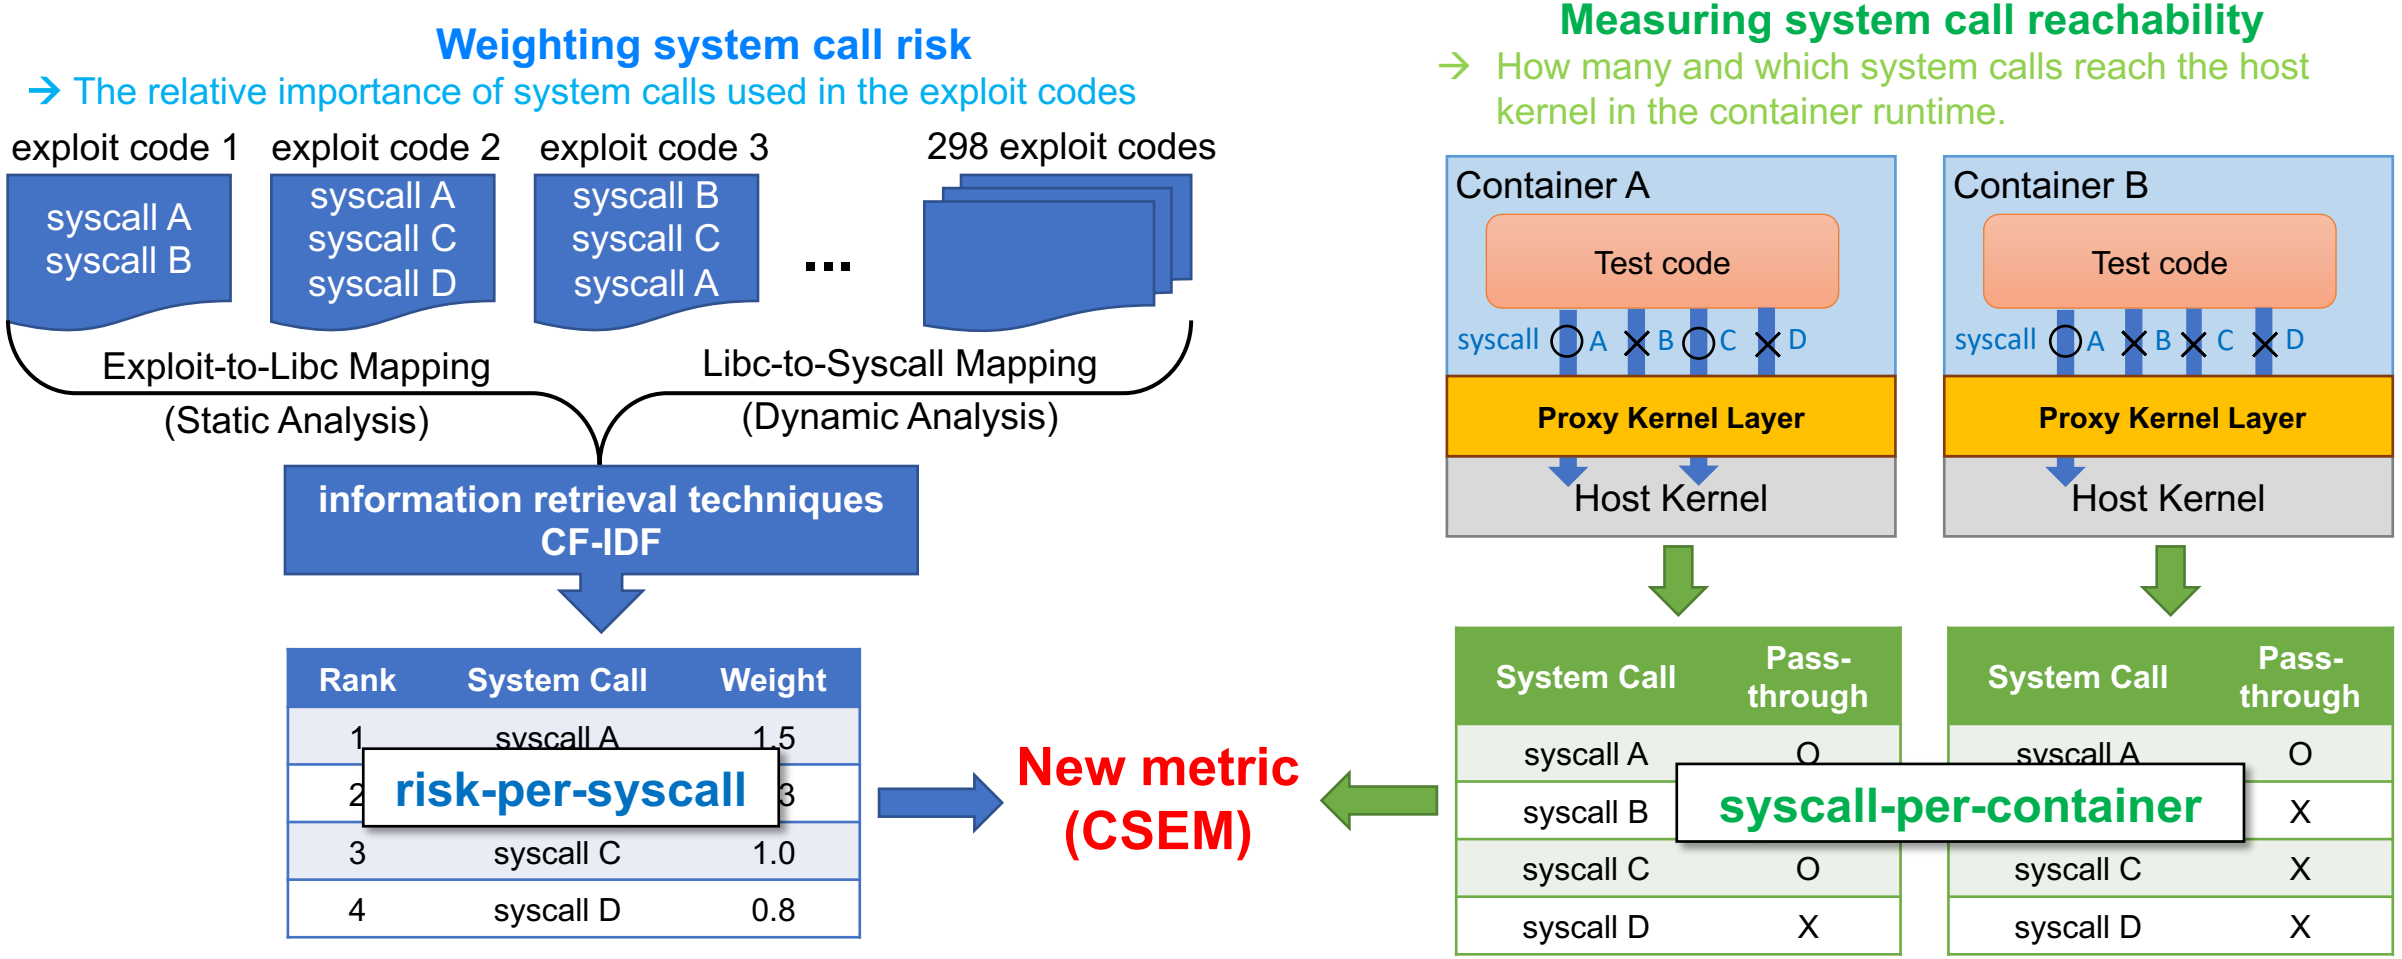

Approach Overview

- An unified metric for comparing container security
  - Two perspectives: **Risk** and **Reachability** of system call

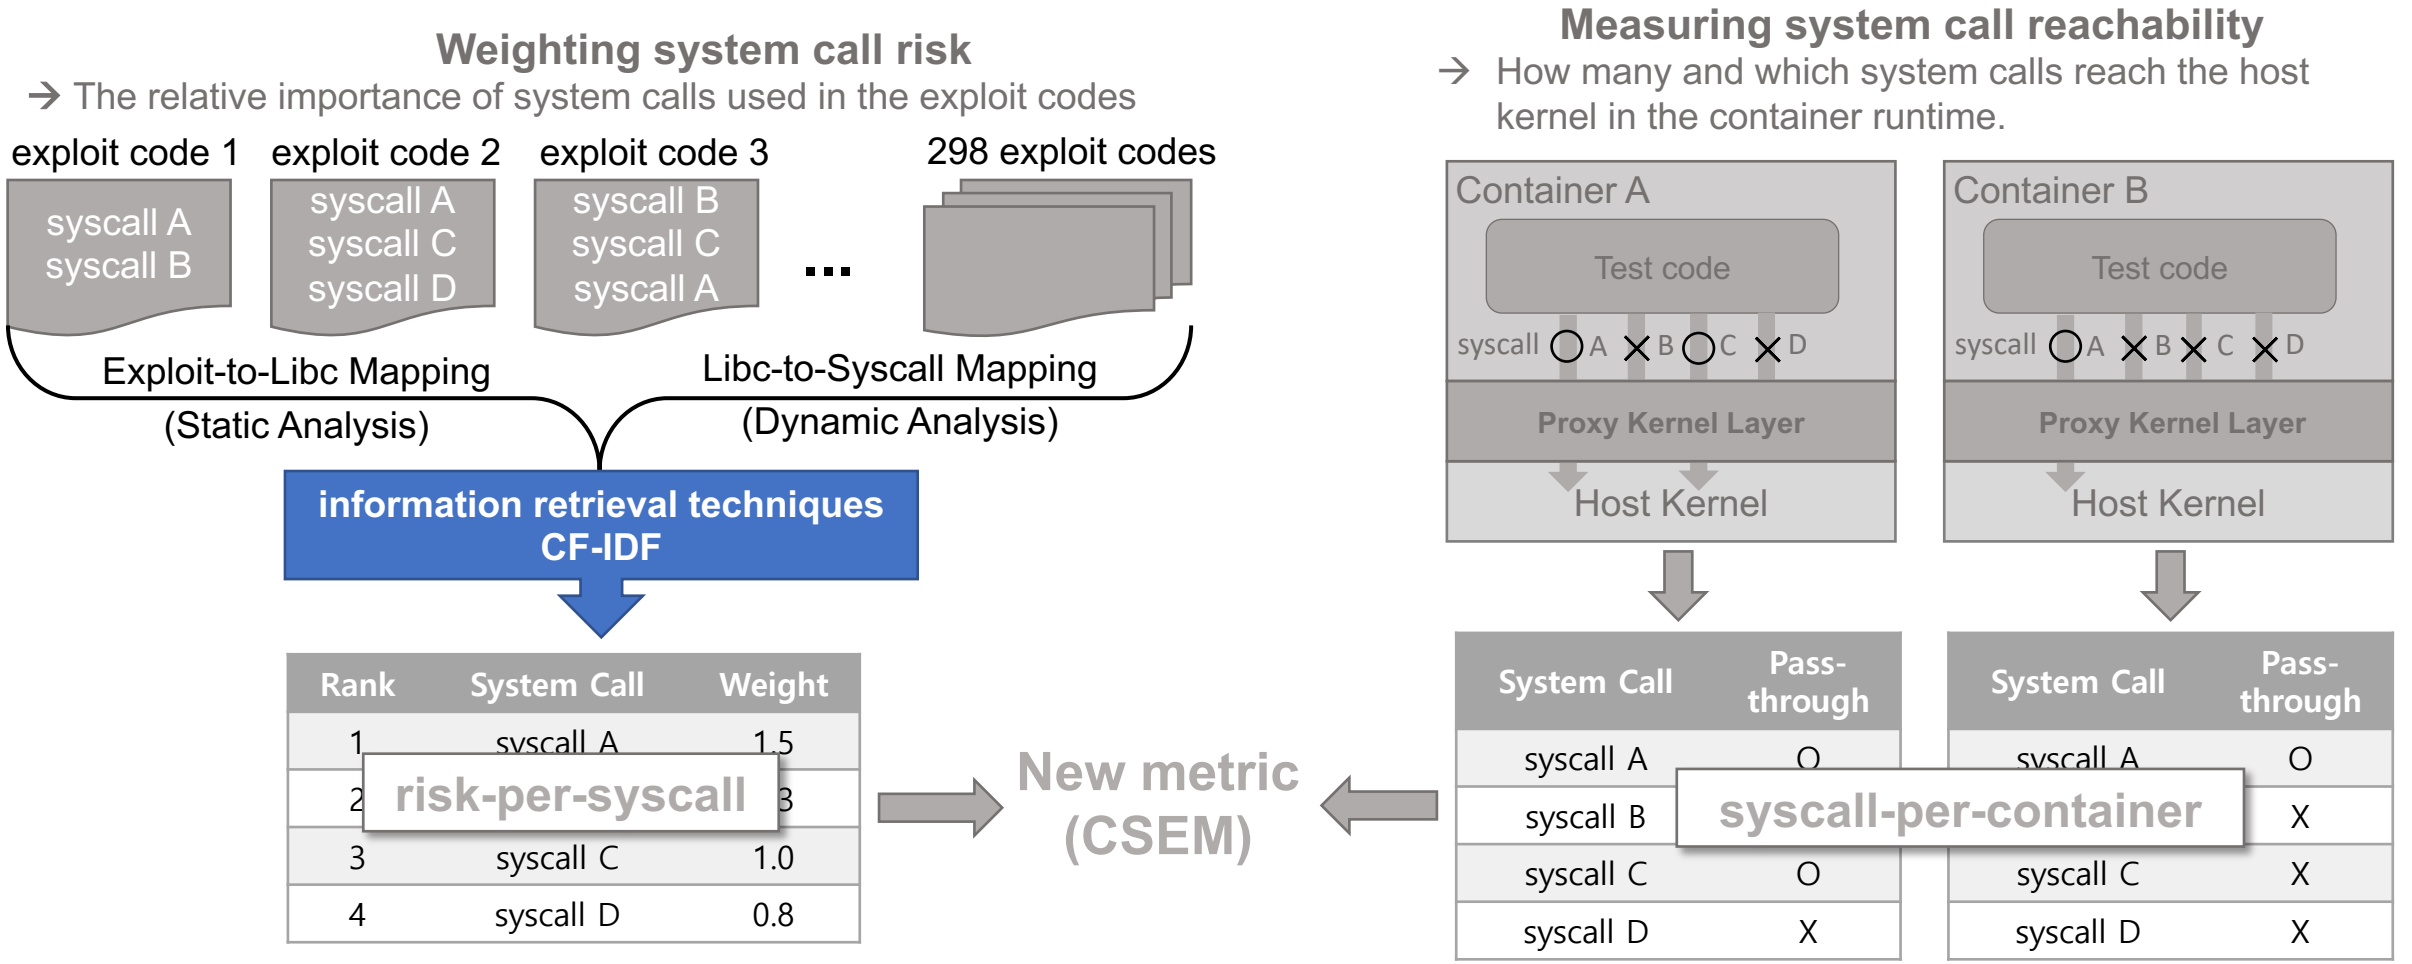

SCAR: System Call Assessment of Risk

- **System call risk weight assignment**: **CF-IDF**
  - Variation of TF(Term Frequency)-IDF(Inverse Document Frequency) for the security quantification

	Term A	Term B	Term C
Doc 1	0	0	1
Doc 2	2	0	0
Doc 3	1	2	1
DF	2	1	2
IDF	Log(3/**2**)	Log(3/**1**)	Log(3/**2**)

TF: Term frequency within a document

DF: Document frequency that include the Term across documents

SCAR: System Call Assessment of Risk

- **System call risk weight assignment**: **CF-IDF**
  - Variation of TF(Term Frequency)-IDF(Inverse Document Frequency) for the security quantification

    ➔ **Term** in **Document** = **System call** in **Exploit code**

	Syscall A	Syscall B	Syscall C
Exploit 1	0	0	1
Exploit 2	2	0	0
Exploit 3	1	2	1
DF	2	1	2
IDF	Log(3/**2**)	Log(3/**1**)	Log(3/**2**)

TF: Frequency in an exploit code may not mean anything.

IDF: Commonly used system calls in exploit codes is less important
e.g., `close`, `brk`, `exit`, `nanosleep`

SCAR: System Call Assessment of Risk

- **System call risk weight assignment**: **CF-IDF**
  - Variation of TF(Term Frequency)-IDF(Inverse Document Frequency) for the security quantification

    ➔ **CF (Class Frequency):** Document Frequency that include the Term within a Class

		Syscall A	Syscall B	Syscall C
Class α	Exploit 1	0	0	1
Class β	Exploit 2	**2**	0	0
	Exploit 3	**1**	2	1
CF of Class β		2	1	1

Syscall	Per-syscall risk weight
Syscall A	0.35
Syscall B	0.48
Syscall C	0.18

CF: The more consistently appearing terms within a class is important to the attack logic of the class

  - Per-syscall risk weight: average CF-IDF values across all exploit codes

Approach Overview

- An unified metric for comparing container security
  - Two perspectives: **Risk** and **Reachability** of system call

Weighting system call risk

→ The relative importance of system calls used in the exploit codes

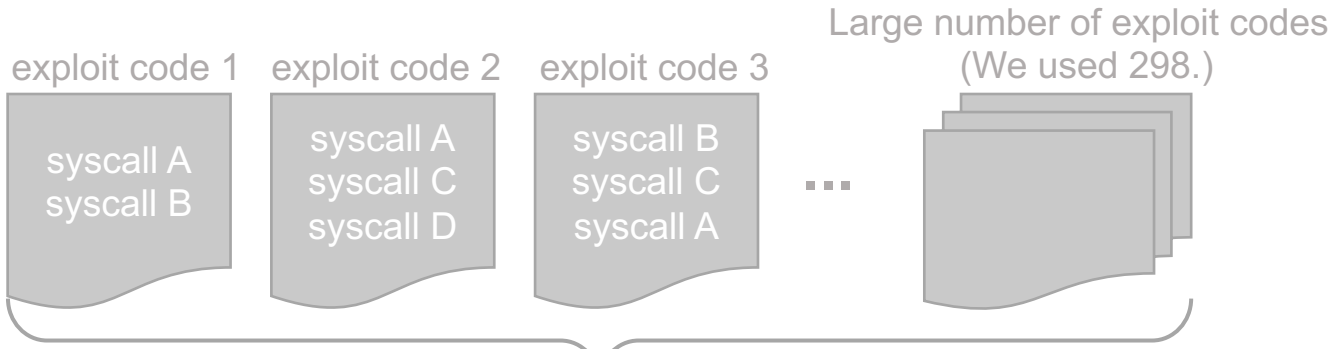

information retrieval techniques

Rank	System Call	Weight
1	syscall A	1.5
2		
3	syscall C	1.0
4	syscall D	0.8

risk-per-syscall

New metric (SecQuant)

Measuring system call reachability

→ How many and which system calls reach the host kernel in the container runtime.

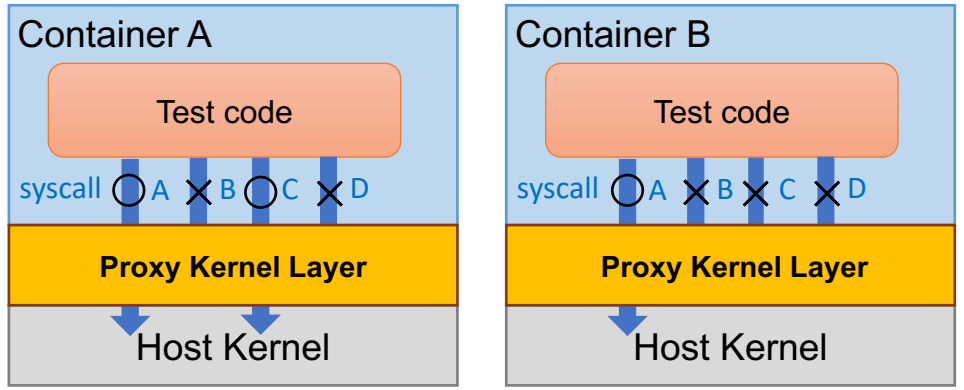

System Call	Pass-through
syscall A	O
syscall B	X
syscall C	O
syscall D	X

System Call	Pass-through
syscall A	O
syscall B	X
syscall C	X
syscall D	X

syscall-per-container

System Call Reachability Test Part

SCED (System Call Exposure Discovery)

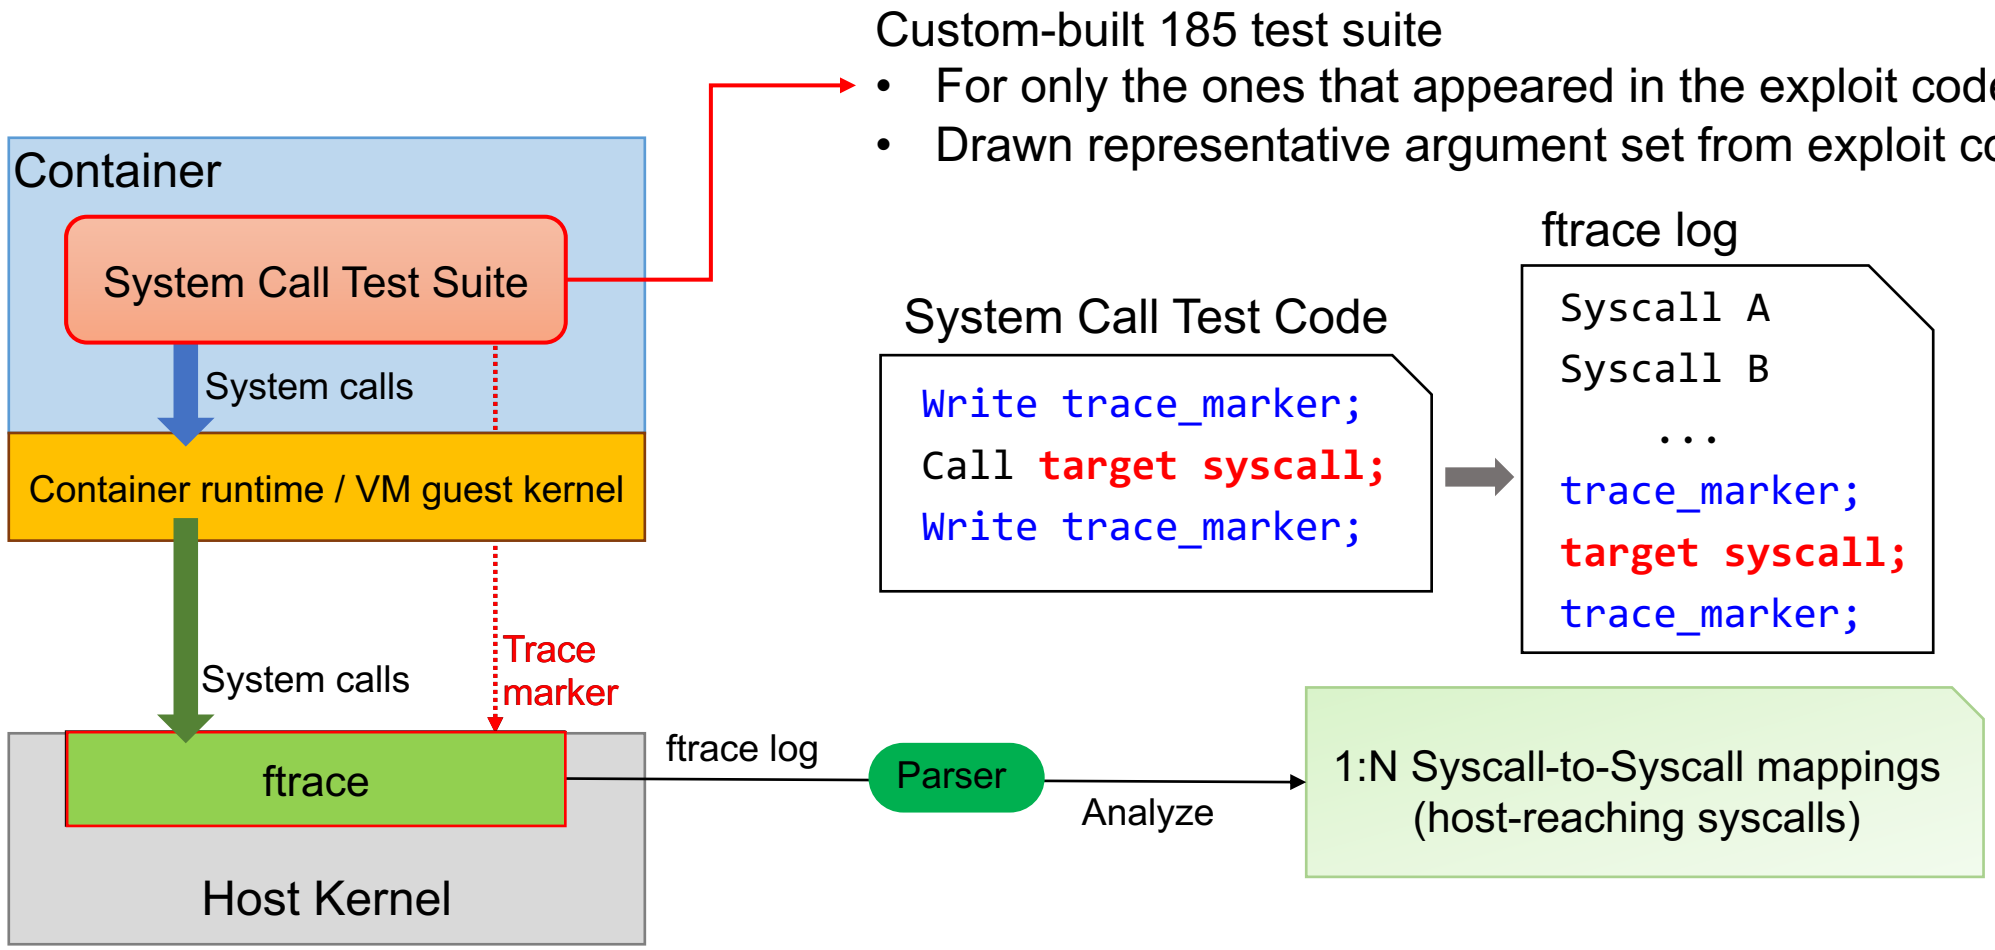

Example: gVisor vs. Kata Containers

- Set of system calls observed at the host kernel
  - Three kinds: Pass-through, Derived, Blocked

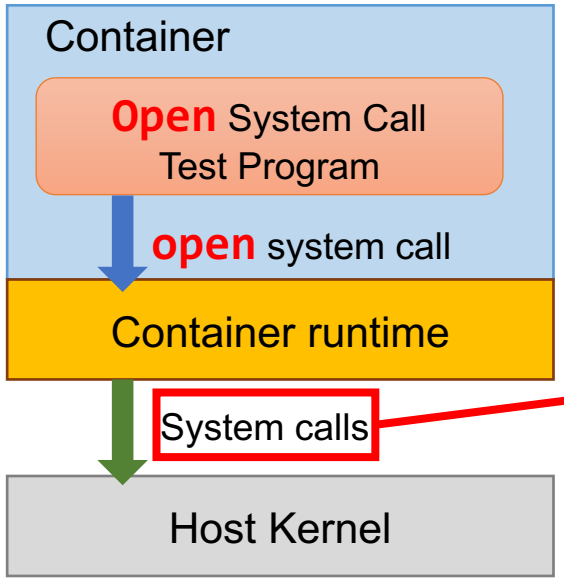

gVisor (runsc-ptrace)		Kata Containers (kata-qemu)	
Host-reaching syscalls	Count	Host-reacing syscalls	Count
futex	20	futex	6
ptrace	10	write	3
openat	4	read	3
newfstat	4	ppoll	3
wait4	2	openat	3
fcntl	2	setresuid	2
sendmsg	1	setresgid	2
recvmsg	1	newfstatat	1
dup	1	geteuid	1
close	1	getegid	1

Derived system calls: A syscall is converted into other set of syscalls

Pass-through *(include equivalent syscalls):* A syscall arrives at the host kernel

CSEM: Container Syscall Exposure Measure

- Combining **SCAR**(risk-weights-per-syscall) and **SCED**(syscall-per-container)

per-syscall risk weight vector

$$CSEM(W, D, r) = \sum_{s} W_s + \sum_{s} \sum_{d \in D_s} \frac{W_d \cdot r}{|D_s|}$$

observed syscall sets

reduction ratio $(0 \leq r \leq 1)$
: control the risk-weight contribution of the derived syscalls

$\sum_{s} W_s$ → **Pass-through syscalls**

the whole risk weight values

reduction ratio

$\sum_{s} \sum_{d \in D_s} \frac{W_d \cdot r}{|D_s|}$ → **Derived syscalls**

Average over the set size of D_s

(D_s: a set of derived system calls from a system call s)

Experimental evaluation: Container Runtime Security Analysis

- Container Syscall Exposure Measure Score Comparison
  - Baseline Container Runtimes: runc, crun, LXC, and sysbox
  - Secure Alternatives: gVisor (runsc-ptrace, runsc-kvm) and Kata containers (kata-qemu. kata-clh)

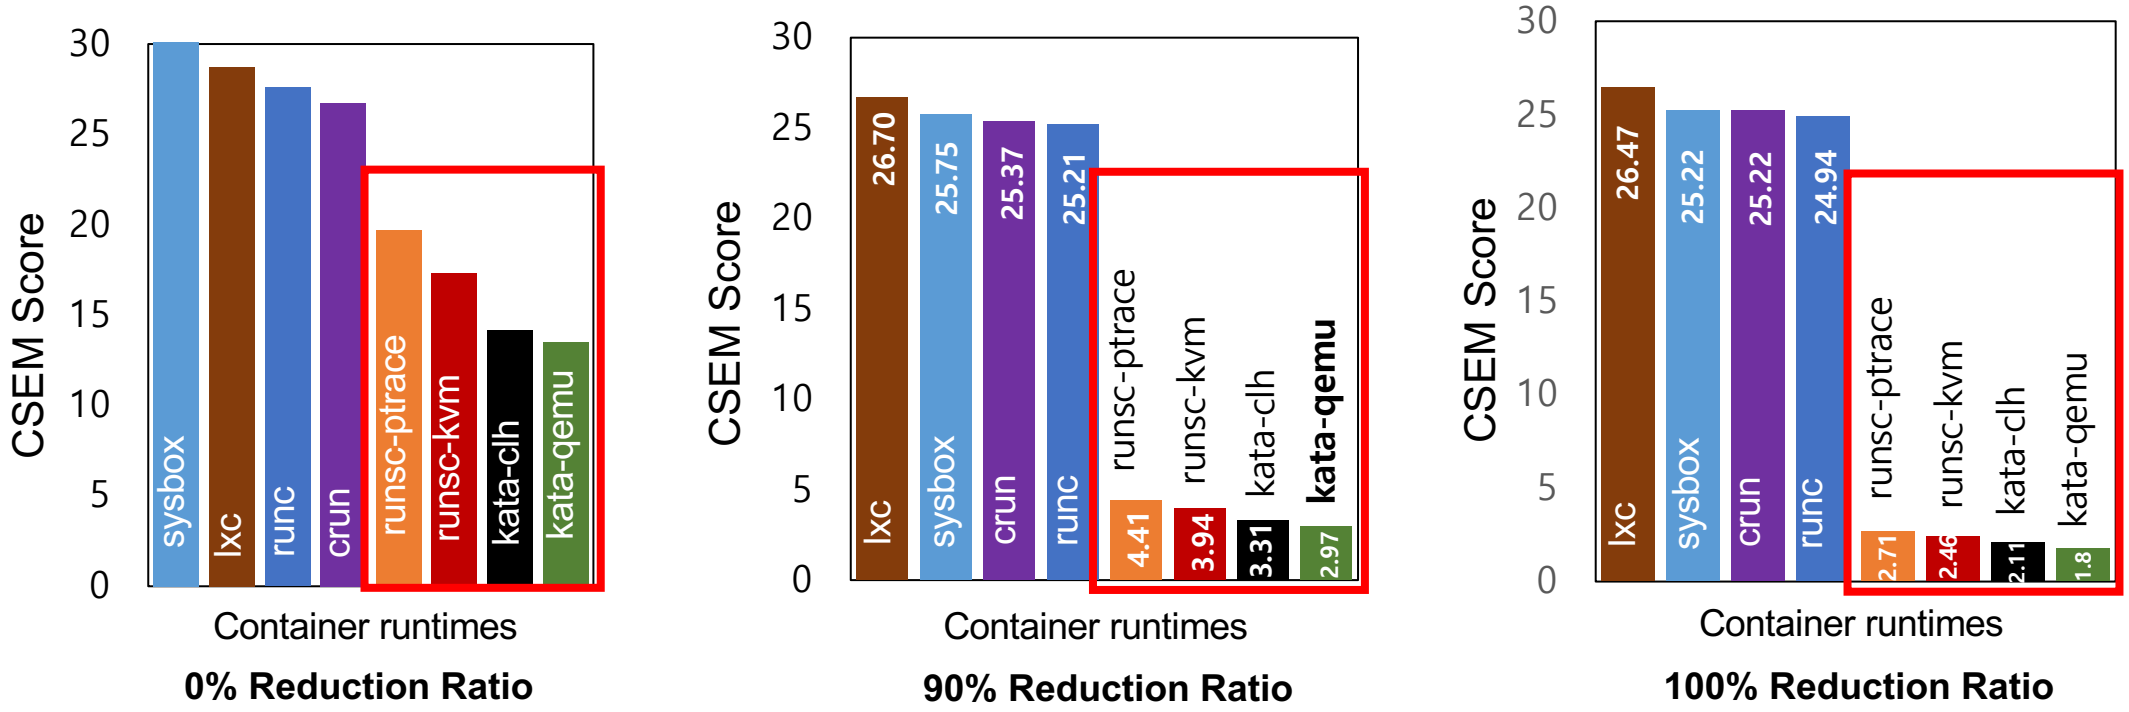

  - Only pass-thru system calls are used in CSEM calculation
  - Derived system calls are treated equally as pass-thru system calls
  - Secure containers have 4.2~7.5 times less syscall exposure than non-security-oriented ones

Experimental evaluation: Container Runtime Security Analysis

- Historical Trends Across Versions
  - Across a 4.5-year history
  - 31, 35 and 22 versions of Docker (runc), gVisor (runsc-ptrace/kvm), and Kata (kata-qemu/clh)

Docker

Event-1
version bump-up for runc, containerd and golang
+ epoll_wait
+ futex

gVisor

Event-2
QEMU upgrade in Kata
- write
- futex

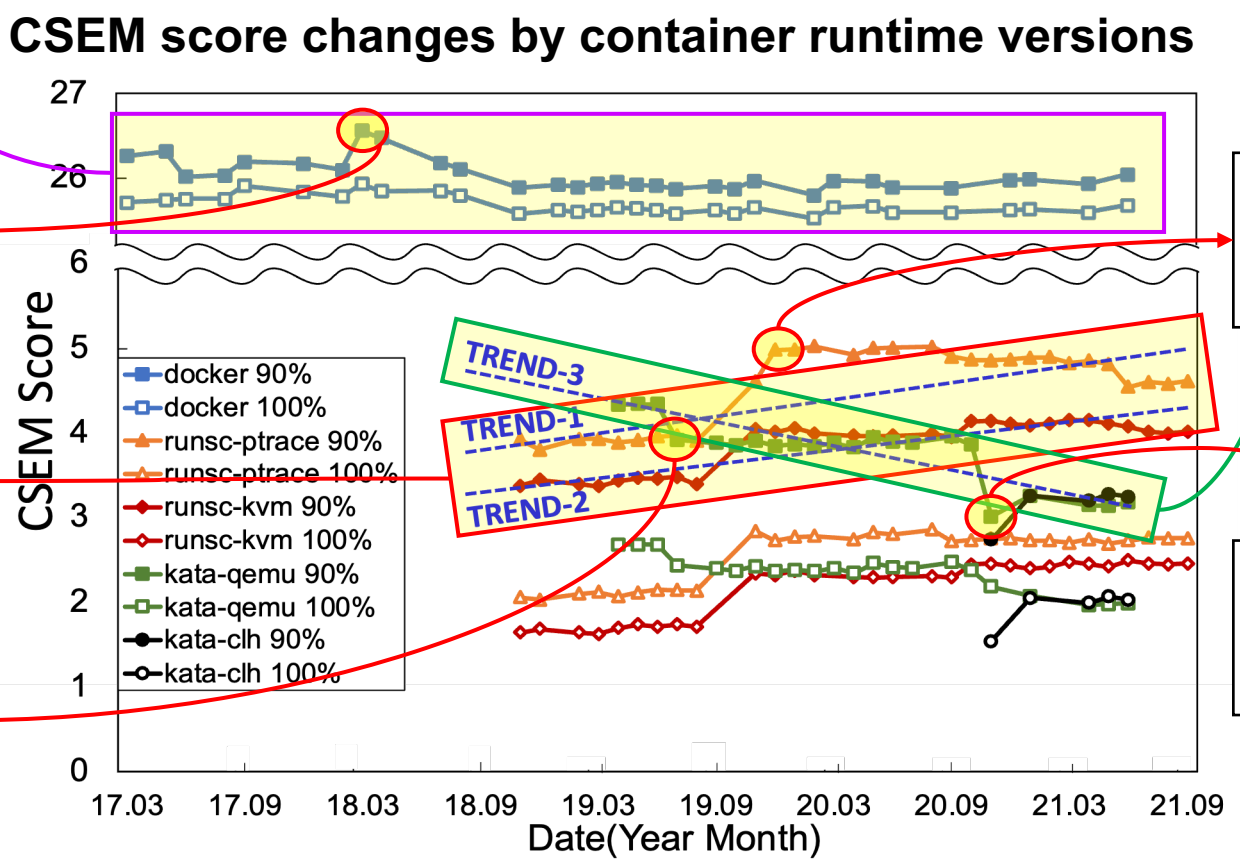

Kata Containers

Event-3
Replace getcpu with golang API
- getcpu

Event-4
- gettimeofday
- clock_gettime

Conclusion

- First attempt at quantifying the security of secure containers.
  - System call risk weighting + System call reachability testing
- Secure containers still allow a significant number of system calls to reach the host kernel
- Future enhancements
  - Exploit code analysis for broader area
  - Handling of vulnerabilities triggered by non-syscalls such as BufferOverflow, Memory Corruption attacks
  - System call weighting refinement by adding benign code analysis
  - System call argument information in tracing
  - Finding proof of successful attacks using:
    - ► Pass-through syscalls with modified arguments
    - ► Derived syscalls

Thank you